

How the Financing of Balance-Sheet Policy Shapes Its Effects*

Dawis Kim[†]

October 5, 2026

[Click here for the latest version](#)

Abstract

I ask whether the financing of Federal Reserve balance-sheet policy shapes its effects on output and the price level. In a monthly U.S. SVAR I identify two balance-sheet shocks, one financed by reserves and one by overnight reverse repurchases (ON RRP). Only reserve financing transmits to output and the price level. A reserve-financed contraction raises the risk premium in corporate bond spreads and lowers both. The reserve-financed shock is the main source of the unexpected movements in output and the price level over the quantitative easing and tightening that reserves financed. Which liability finances a balance-sheet policy shapes what that policy does.

Keywords: quantitative tightening, quantitative easing, central bank liabilities, overnight reverse repurchases, excess bond premium, sign and zero restrictions, shock heteroskedasticity

JEL codes: C11, C32, E51, E52, E58, G21

*I am deeply grateful to my advisors, Eric M. Leeper, Eric R. Young, and Leland E. Farmer, and I owe special thanks to Tao Zha. For comments at the Federal Reserve Bank of Atlanta I thank R. Anton Braun, Jonathan L. Willis, Salomé Baslandze, Larry D. Wall, and all seminar participants there. For comments at the University of Virginia, Duke University, and Federal Reserve Bank of Richmond Joint Research Workshop I thank Huberto M. Ennis and Alexander L. Wolman. For comments at the University of Virginia practice job talk I thank Andrew Glover and everyone else who commented. All errors are my own.

[†]University of Virginia. cbe2zn@virginia.edu.

1 Introduction

Do quantitative easing (QE) and quantitative tightening (QT) move output and the price level? This paper finds that the answer depends on which liability finances them. A balance-sheet operation moves the securities of the System Open Market Account (SOMA). The liability side moves by the same amount, in reserves or in overnight reverse repurchases (ON RRP). A liability finances the operation when the balance comes to rest in it. The Federal Reserve (Fed) pays for every purchase with reserves, and the balance then stays there or moves into ON RRP as the holders of the cash decide. The liability that finances the operation determines who receives the Fed's interest, banks on reserves or money market funds (MMFs) on ON RRP.




Figure 1: Liability transition across balance-sheet policy.

Note. Bars show the change in reserves and in ON RRP over the episode windows of Table A.1 as a percent of each episode's total asset change. The number above each bar is that share. QE episodes point upward and QT episodes downward. Table A.1 reports the changes in USD billions. The two shares do not sum to 100 percent, and the remainder is the change in the TGA, currency, and the residual liabilities, which is net positive in every episode. Reserves fell by more than total assets over QT1. Inside the COVID QE window the ON RRP buildup broadly matched the contemporaneous TGA drawdown, as Section 2 documents. Source: Federal Reserve H.4.1 and NY Fed SOMA holdings.

Since 2008 the Fed has run two rounds of QE and two of QT. The liability that financed them

changed from one round to the next. Reserves financed the 2008 to 2014 purchases, and QT1 drained them. Both liabilities financed the COVID purchases. QT2 left the reserves nearly intact and drained the ON RRP balances of MMFs instead (Figure 1).

The two financings differ in how they could reach the cost of credit for firms. Under reserve financing, banks earn the margin of interest on reserves over their funding costs, and that income would add to bank capital. Reserves would also relax the leverage constraint of banks when they replace assets that count against it. Either channel would move the risk premium banks require to bear credit risk. MMFs pass the ON RRP rate to their shareholders as yield and keep none of it as capital. Unlike banks, MMFs choose where the balance sits and move it to whichever of ON RRP and Treasury bills pays more. If the transmission differs with the liability, the difference should appear first in the risk premium that moves the cost of credit for firms.

The two financings motivate two balance-sheet shocks, a reserve-financed shock (RFBS) and an ON RRP-financed shock (OFBS). I identify them in an eight-variable monthly U.S. SVAR, alongside demand, supply, monetary policy (MP), and excess bond premium (EBP) shocks. The EBP is a risk premium in corporate bond spreads beyond the compensation for expected default. A reserve-financed contraction raises the EBP and lowers GDP and the GDP deflator. An ON RRP-financed contraction moves neither the EBP nor GDP but lowers the deflator. Once the model controls for the arbitrage of ON RRP holders between short-term assets, the decline disappears.

Holding the EBP fixed removes both declines after a reserve-financed contraction. Following [Brunnermeier et al. \(2021\)](#), I combine the RFBS shock with a sequence of EBP shocks that keep the EBP response at zero. With the EBP held at zero, the declines in GDP and the deflator are no longer present. The counterfactual attributes the effects of a reserve-financed contraction on GDP and the deflator to the rise in the EBP.

The shares of both balance-sheet shocks in the forecast variance of GDP and the deflator grow with the horizon, to about a tenth at five years. The RFBS share of the EBP grows faster and reaches 18% within two years, close to the share of the EBP shock itself.

Each balance-sheet shock is the main source of the unexpected movements in GDP and the

deflator when its own liability mainly financed the balance sheet. Reserves mainly financed QE and QT1, and over both the RFBS shock is the largest contributor to the GDP, deflator, and EBP forecast errors. It pushes GDP and the deflator above their forecasts over QE and below over QT1, and it moves the EBP the other way. ON RRP mainly financed the drain of QT2. Over the drain the OFBS shock is the largest contributor to the GDP and deflator forecast errors and pulls both below their forecasts. Once the model controls for that arbitrage, the OFBS contribution to the deflator over the drain falls to near zero. The RFBS shock remains the largest contributor over QE and QT1.

The volatility of the shocks shifts across the sample. The two decades since 2008 contain the financial crisis, the COVID collapse, and the post-2021 inflation surge, and the months around COVID are extreme within their phase. I model both with the shock volatility and the heavy-tailed shocks of [Brunnermeier et al. \(2021\)](#). Sign, zero, and narrative restrictions give the shocks their names. The zero restrictions carry the timing assumption of [Leeper et al. \(1996\)](#) and [Christiano et al. \(2005\)](#). Real activity and prices do not respond to a balance-sheet shock within the month. The balance sheet changes gradually after an announcement ([Nakashima et al., 2024](#)), and its real effects arrive after the quantities accumulate.

No existing sampler imposes zero restrictions on a heteroskedastic model. This paper develops the heteroskedastic absorb-and-rotate sampler with zero restrictions (HARS-Z). HARS-Z extends the heteroskedastic absorb-and-rotate sampler (HARS) of [Kim and Zha \(2026\)](#), which combines the heteroskedastic likelihood of [Brunnermeier et al. \(2021\)](#) with sign restrictions and with the narrative restrictions of [Antolín-Díaz and Rubio-Ramírez \(2018\)](#).

To my knowledge, no existing paper estimates the real effects of balance-sheet policy separately by the liability that finances it. Most of the empirical literature measures the financial-market effects of balance-sheet policy in event studies around announcements, for QE ([Krishnamurthy and Vissing-Jorgensen, 2012](#); [Swanson, 2021](#)) and for QT ([Smith and Valcarcel, 2023](#); [Lloyd and Ostry, 2024](#); [Lu and Valcarcel, 2024](#)). Theoretical work models the transmission of balance-sheet policy ([Chen et al., 2012](#); [d'Avernas et al., 2024](#); [Kumhof and Salgado-Moreno, 2024](#); [Adrian et al., 2025](#)). In these models reserves, or the short-term debt of a consolidated government, finance the

purchases. [Nakashima et al. \(2024\)](#) estimate the real effects of the Bank of Japan’s balance-sheet policy with a different identification. Their shocks separate the size of the balance sheet from its asset composition, not the liability side.

The rest of the paper proceeds as follows. [Section 2](#) documents the balance-sheet evidence. [Section 3](#) presents the empirical model, the data, and the identification. [Section 4](#) develops HARS-Z and establishes its validity. [Section 5](#) reports the impulse responses, their robustness, and the decompositions. [Section 6](#) concludes. The appendices contain the balance-sheet data, the estimation diagnostics, the formal statement of the narrative restrictions, the review of HARS and the theory of HARS-Z, the proofs, and the robustness exercises.


2 Two Ways of Financing the Balance Sheet

2.1 The Two Sides of the Balance Sheet


QE expands the Fed’s balance sheet by buying securities, and QT shrinks it by letting securities run off at maturity without reinvestment. [Figure 2](#) plots the assets and liabilities. Every QE raises both sides and every QT lowers them together. The balance-sheet identity fixes the totals and leaves open which liabilities absorb the change.

The asset side. The four episodes differ in size and not in composition. Both expansions raised Treasury and MBS holdings together, and both contractions lowered them together ([Figure 2a](#)). The COVID QE raised total assets by \$4.3 trillion, roughly twice the \$2.3 trillion of the 2008 to 2014 expansion. QT2 ran longer and larger than QT1, with coupon runoff of \$1.45 trillion against \$0.38 trillion ([Appendix A](#)).

The liability side. The episodes diverge in which liability they moved. The Fed buys a security from a dealer and pays the dealer’s bank in reserves. The bank credits the dealer with a deposit. The reserves stay where they are unless the deposit moves into an MMF that places the cash in the



(a) Fed asset holdings.



(b) Fed liabilities.

Figure 2: The Fed's balance sheet.

Note. Panel (a) plots the holdings of Treasury coupons, Treasury bills, and MBS, and panel (b) plots the reserve, ON RRP, and TGA balances, in USD billions. Each line is a monthly average, and the shaded spans are the episode windows of Table A.1. Source: Federal Reserve H.4.1.

ON RRP facility, which turns the reserves into ON RRP.¹ When the Fed lets a security run off, the Treasury repays it from its account at the Fed and refills that account by issuing new debt. Reserves fall when banks or their depositors buy the new debt, and ON RRP falls when MMFs buy it with cash they take out of the facility. Whether the Fed buys or lets a security run off, the holders of the cash decide which liability carries the balance (link 1 of Figure 3). They decide within the terms the

¹Suppose the Fed buys \$100 of securities from a dealer. The dealer's bank receives \$100 of reserves, and the dealer receives a \$100 deposit. The dealer had borrowed from an MMF in repo to hold the security, and it uses the deposit to repay the MMF. The \$100 of reserves now sit at the MMF's bank instead of the dealer's bank. If the MMF places the \$100 at the ON RRP facility, the reserves leave the banking system and become ON RRP. If the MMF buys a bill from another investor instead, the reserves stay in the banking system.

Fed sets for the facility.

Each share below is percent of the episode's total asset change. Reserves absorbed 78% of the 2008 to 2014 expansion. QT1 reversed less than a third of that expansion and drained essentially the entire contraction from reserves. The COVID QE split between reserves (46%) and ON RRP (37%). QT2 drew 92% of the contraction from ON RRP and 13% from reserves (Appendix A).

2.2 The ON RRP Facility from Buildup to Drain

The ON RRP buildup. The ON RRP balance was \$154 billion when QT1 began and near \$2 trillion when QT2 began. The COVID purchases settled in reserves at first, and reserves kept rising into late 2021 as the buying continued (Figure 2b). The balance grew after March 2021, when the Treasury cut bill issuance and drew down the Treasury General Account (TGA) and the released cash flowed through MMFs into ON RRP (Wong, 2021). The expiry of the temporary Supplementary Leverage Ratio (SLR) relief in March 2021 added to it.²

The QT2 drain and the ON RRP facility. QT2 drained the ON RRP facility where QT1 had drained reserves. The Committee's 2022 principles state that the runoff ends when reserves reach the level it judges consistent with an ample reserves regime (Federal Open Market Committee, 2022b). The drain ran from mid-2023, a year after the runoff began. Heavy Treasury bill issuance pushed T-bill yields above the ON RRP rate, and MMFs pulled their cash out of the facility and bought bills. Government funds left first, and their share of the balance fell from 80 percent in 2023:M7 to 35 percent in 2025:M9 (Figure A.1). The shift lowered ON RRP balances and left reserves intact. Over the first three years of the runoff, securities holdings fell by \$1.3 trillion while reserves rose by \$350 billion (Logan, 2024).

Members of the Federal Open Market Committee watched the ON RRP facility as they set the

²The SLR requires large banks to hold tier 1 capital against total leverage exposure without risk weights. In April 2020 the Fed temporarily excluded Treasuries and reserve balances from the exposure measure to ease strains in the Treasury market. The exclusion expired as scheduled on March 31, 2021, and reserves and Treasuries again counted toward bank leverage. Banks discouraged large deposits, the cash moved to MMFs, and they found few bills to buy (Afonso et al., 2025). Wong (2021) attributes the largest share of the ON RRP buildup to the fall in bill supply during the TGA drawdown, with the SLR expiry and the June 2021 rate adjustment adding to it.

pace of the runoff. [Waller \(2024\)](#) read the cash remaining in the facility as excess liquidity and as room to reduce securities holdings further. [Logan \(2024\)](#) attributed the decline in the facility to Treasury bill issuance and warned that further runoff would fall on reserves once the facility reached a low level. The Committee saw the same risk. In June 2024 it lowered the monthly Treasury redemption cap from \$60 billion to \$25 billion so that the runoff would not fall sharply on reserves ([Federal Open Market Committee, 2024](#)).

The ON RRP claim also moves duration. When reserves finance a purchase, the Fed takes duration out of private hands and long-term yields fall ([Krishnamurthy and Vissing-Jorgensen, 2012](#); [Chen et al., 2012](#)). When it lets those securities run off, private investors take the duration back ([Crawley et al., 2022](#); [Wei, 2022](#)). A reader may take the ON RRP-financed episodes differently. In that reading the buildup is MMFs placing cash in the facility, the drain is MMFs exchanging repo with the Fed for bills, and no duration moves either way.

Duration moved in both ON RRP episodes. The COVID purchases took coupons and MBS out of private hands whether the balance settled in reserves or moved into ON RRP. Over QT2 the Fed let \$1.45 trillion of Treasury coupons run off (Table [A.1](#)). Private investors hold those coupons instead, as they hold the coupons that ran off under QT1. Table [A.2](#) shows who took them on. MMFs added Treasury securities and repo lending over QT2, and hedge funds borrowed in repo to hold more Treasuries. Both financings move duration the same way, and Section [3](#) imposes the same term-spread restriction on both balance-sheet shocks for that reason.

2.3 Why the Financing Should Matter

Each financing pays a different holder. The Fed pays interest on both claims, and the two financings differ in who receives it and in what that interest does. Only banks can hold reserves, and interest on reserves goes to them (link 2 of Figure [3](#)). A reserve-financed operation changes the quantity of reserves on bank balance sheets. Banks do not choose that quantity. The Fed sets it, and reserves cannot leave the banking system.




Figure 3: Two ways to finance the same operation, and what each changes for its holder.

Note. A stylized Fed balance sheet. A purchase of securities on the asset side creates reserves, and the balance stays in reserves or moves into ON RRP as the holders of the cash decide. MBS are agency mortgage-backed securities. The third column names who holds each claim, and the fourth names what the claim changes for its holder. The fifth names where the reserve-financed chain reaches the cost of credit for firms. The circled numbers mark the four links the text follows. Interest on reserves is bank income and adds to bank capital when the reserves add to bank balance sheets. In the model of [Sims and Wu \(2021\)](#), reserves that replace assets counting against the leverage constraint of banks relax that constraint. Both routes move the risk premium banks require to bear credit risk. The ON RRP rate passes through money market funds to their shareholders as yield and builds no capital on the way.

Banks earn the rate on reserves and pay the cost of the deposits or wholesale borrowing that fund them, and the margin is bank income.³ That income adds to bank capital, and capital limits how much credit risk, the risk that borrowers do not repay, a bank can carry (link 3). Reserves can also move the risk premium when they do not add to bank balance sheets. In the model of [Sims and Wu \(2021\)](#), reserves that replace assets counting against the leverage constraint of banks relax that constraint (link 4). A looser constraint lowers the risk premium banks require.

The interest on ON RRP goes almost entirely to MMFs, mostly to government funds.⁴ MMFs keep none of that interest as capital and pass it to their shareholders as yield. Unlike banks, MMFs choose where the balance sits and move it to whichever of the facility and bills pays more. Moving

³When the European Central Bank raised its deposit facility rate in 2022, euro area banks with more reserves gained net interest income and book equity relative to other banks ([Fricke et al., 2026](#)).

⁴At the December 2022 peak MMFs held 92 percent of the ON RRP balance, and government funds held 77 percent. Government funds must invest at least 99.5 percent of their total assets in cash, government securities, and fully collateralized repo ([U.S. Securities and Exchange Commission, 2026](#)). Commercial paper, the short-term unsecured debt that firms issue, falls outside that set. Prime funds, which can buy commercial paper, never held more than 17 percent (Figure A.1).

the balance between the two changes what MMFs hold, not what they earn.⁵ The yield reaches households and firms as interest income, of the same kind that depositors receive under reserve financing, and it builds no capital.⁶ No capital changes along the ON RRP-financed chain, and the credit risk that banks can carry stays where it was.

The reserve-financed chain ends at bank capital and at the leverage constraint of banks, and the ON RRP-financed chain ends at the shareholders of MMFs. Intermediaries with more capital require a smaller risk premium to lend to the same borrower, and those with less capital require a larger one (He and Krishnamurthy, 2013).⁷ If the two chains end in different places, the difference should appear first in the cost of credit for firms. Section 5 asks whether the difference shows in the risk premium and in GDP and the price level.

The terms and the realized split. The Fed set the terms that govern the split between reserves and ON RRP. It set the operating framework, the ON RRP award rate, the per-counterparty caps, and the runoff caps (Federal Open Market Committee, 2022b,a). Within those terms, Treasury issuance, bank regulation, and the portfolio choices of MMFs affected the realized split, as Section 2.2 records (Afonso et al., 2025). Those terms set the financing. If the transmission differs with the financing, they also shape what an operation of a given size does to GDP and the price level. This paper does not model how the Fed chose those terms. I take the realized financing as given and identify two balance-sheet shocks, one financed by reserves and one financed by ON RRP.

The gap in existing work. Existing work prices balance-sheet policy without separating the liability that finances it. Crawley et al. (2022) in an FRB/US simulation and Wei (2022) in a

⁵Government MMFs exchanged their ON RRP balance for other assets. From 2023:M5 to 2024:M5 they reduced their ON RRP holdings by \$1.39 trillion and added \$1.36 trillion of Treasury securities and \$0.58 trillion of repo with private counterparties. Their total assets rose by \$0.44 trillion.

⁶In the Financial Accounts, households and nonprofit organizations held 57 percent of money market fund shares at the end of 2022, and nonfinancial businesses held 17 percent (table F.521). The table does not split fund types, and government funds held 78 percent of fund assets on that date (table S.123) (Board of Governors of the Federal Reserve System, 2026).

⁷Twenty-one of the twenty-six primary dealers are bank-affiliated (Federal Reserve Bank of New York, 2026). Bank capital limits the corporate bonds these dealers can hold. Banks are also the largest lenders to firms, and firms move from loans to bonds when banks tighten their terms (Becker and Ivashina, 2014). Banks with less capital push borrowers into the bond market and raise the spread there.

preferred-habitat model convert a runoff into an equivalent policy-rate increase. Both price it through the duration risk that private investors absorb. [Wei \(2022\)](#) traces which liability a runoff drains, reserves when banks buy the reissued coupon securities and ON RRP when MMFs buy bills. He does not estimate the two cases. In [Kumhof and Salgado-Moreno \(2024\)](#) and [d'Avernas et al. \(2024\)](#) reserves finance balance-sheet policy and banks transmit it, and neither model includes an ON RRP facility.

3 Empirical Model

Let y_t be the $n \times 1$ vector of endogenous variables and x_t the $r \times 1$ vector of predetermined terms, including the constant and p lags of y_t . The structural equation is

$$y_t' A_0 = x_t' A_+ + \eta_t', \quad (1)$$

where A_0 is the $n \times n$ nonsingular structural impact matrix and A_+ is the $r \times n$ matrix stacking the constant and lag coefficients. The vector η_t collects the model shocks, independent across equations and time conditional on the regime.⁸

Following the deterministic variance-regime specification of [Brunnermeier et al. \(2021\)](#), let

$$s_t : \{1, \dots, T\} \rightarrow \{1, \dots, h\}$$

be a known date-to-regime map that assigns each date to one of h volatility regimes. The shocks are heteroskedastic across regimes and heavy-tailed within them,

$$\eta_{i,t} \mid \nu_{i,t}, s_t = k \sim \mathcal{N}(0, \lambda_{i,k} \nu_{i,t}), \quad i = 1, \dots, n, \quad (2)$$

where $\lambda_{i,k} > 0$ is the relative variance of shock i in regime k and $\nu_{i,t}$ is a latent scale with an

⁸Under this right-multiplication convention, which (1) shares with [Kim and Zha \(2026\)](#) and [Arias et al. \(2018\)](#), columns of A_0 correspond to equations.

inverse-gamma distribution, as in [Brunnermeier et al. \(2021\)](#). The inverse-gamma parameters make each standardized shock a Student- t with unit scale, and the variance of shock i in regime k is proportional to $\lambda_{i,k}$. I normalize the regime-specific variances to satisfy

$$\frac{1}{h} \sum_{k=1}^h \lambda_{i,k} = 1, \quad i = 1, \dots, n,$$

which fixes the scale of each shock variance path across regimes.

Post-multiplying (1) by A_0^{-1} gives the reduced form

$$y'_t = x'_t B + u'_t, \quad B = A_+ A_0^{-1}, \quad u'_t = \eta'_t A_0^{-1}, \quad (3)$$

where the regime-specific reduced-form covariances

$$\Sigma_k = \text{Var}(u_t \mid s_t = k) = (A_0^{-1})' \Lambda_k A_0^{-1}, \quad \Lambda_k := \text{diag}(\lambda_{1,k}, \dots, \lambda_{n,k}), \quad (4)$$

identify the columns of A_0 up to signed permutation whenever the variance ratios across regimes are not mutually proportional ([Rigobon, 2003](#); [Lanne et al., 2010](#)). The covariances identify (A_0, Λ) but do not label the shocks. In this paper, the economic labeling comes from the restricted rotation, and the two roles stay separate throughout.

3.1 Data and Priors

I estimate the model of [Section 4](#) on a monthly U.S. system with eight variables ($n = 8$) and $p = 6$ lags. The sample is November 2008 through June 2026.⁹

The system contains eight variables. The first four are the one-year Treasury yield, the log of SOMA holdings, $\log(\text{Reserves} + \text{ON RRP})$, and the ON RRP ratio $\text{ON RRP}/(\text{Reserves} + \text{ON RRP})$. SOMA holdings cover Treasury notes and bonds plus agency MBS and exclude bills.¹⁰ The last

⁹The November 2008 start coincides with the financial-crisis interventions that preceded QE1. It drops the pre-crisis regime, where reserves stayed near zero and reserve variance was orders of magnitude smaller.

¹⁰The Fed treats bills separately from the programs. It purchased bills from October 2019 into 2020 to keep reserves

four are the log of real GDP, the log of the GDP deflator (GDPDEF), the term spread, and the excess bond premium (EBP). The term spread is the 10-year Treasury yield less the 3-month bill rate. I map quarterly GDP and the deflator to monthly frequency by Chow–Lin interpolation and aggregate every series to monthly means.

The one-year Treasury yield stands in for the policy rate. At the zero lower bound the funds rate carries no innovation. [Gertler and Karadi \(2015\)](#) take a short-maturity government yield as the policy indicator, on the ground that announcements about future policy operate over a horizon of about two years. Their monthly baseline uses the one-year.

I include the EBP for the reason [Section 2](#) gives. The two financings differ in how they move the risk premium intermediaries require to bear credit risk. The EBP measures that premium in the corporate bond market. It is the part of corporate bond spreads not attributable to expected default risk. [Gilchrist and Zakrajšek \(2012\)](#) read a rise in the EBP as a fall in the financial sector’s capacity to bear that risk.¹¹ In their estimates, a rise in the EBP that is orthogonal to the state of the economy lowers economic activity. [Glover and Oliyide \(2024\)](#) read the EBP as the willingness of lenders to lend beyond what the riskiness of the loans explains. [Gertler and Karadi \(2015\)](#) include it in a system like this one.

[Section 2](#) implies a prediction for the EBP. Suppose the reserve-financed chain adds to the income and capital of banks and the ON RRP-financed chain does not. Intermediaries with more capital require a smaller risk premium, and a reserve-financed expansion would lower the EBP while a reserve-financed contraction would raise it. An ON RRP-financed operation would move the EBP in neither direction. The EBP carries no restriction under either balance-sheet shock, and the data decide whether that pattern appears. [Section 5.1](#) reports how each balance-sheet shock moves it.

Two liability variables enter the system. The size variable $\log(\text{Reserves} + \text{ON RRP})$ sums the two liabilities that absorbed the QE and QT episodes of [Section 2](#). It treats reserves and

ample. It described those purchases as technical measures that do not represent a change in the stance of monetary policy ([Federal Open Market Committee, 2019](#)). Under the 2022 runoff the decline in Treasury holdings covers coupon securities, and bills fall only to the extent that coupon maturities run below the monthly cap ([Federal Open Market Committee, 2022a](#)). Bill holdings track the maturity schedule and the reserve-management purchases rather than the purchase and runoff decisions.

¹¹The part of the spread they remove moves with the borrowers, and what remains moves with the lenders.

ON RRP as close substitutes, two forms of overnight central-bank money that commercial banks and MMFs hold, respectively (Diamond et al., 2024; Copeland et al., 2021). The ON RRP ratio $\text{ON RRP}/(\text{Reserves} + \text{ON RRP})$ measures the sector split of that total. The ratio distinguishes the two contractions. A reserve drain and an ON RRP drain reduce the size variable by the same amount, but they move the ratio in opposite directions. The tables abbreviate the two variables as $\log(R + \text{RRP})$ and $\text{ON RRP}/(R + \text{RRP})$.¹²

The priors follow Sims and Zha (1998) and Brunnermeier et al. (2021). Each element of A_0 has an independent Gaussian prior with mean 100 on the diagonal, 0 off the diagonal, and standard deviation 200. The centering at $A_0 = 100I$ scales each equation to an expected prior standard deviation of 0.01, and the standard deviation of 200 keeps the prior diffuse. A Dirichlet($\alpha = 2$) prior on the normalized variance shares centers the prior on equal variances across regimes. It enforces the normalization that the relative variances of each shock average to one.

Sims–Zha dummy observations center each equation on an independent random walk, except the ON RRP ratio and the term spread, which I center on white noise.¹³ Following Brunnermeier et al. (2021), the latent scales are i.i.d. inverse-gamma and make each standardized shock Student- t with 5.84 degrees of freedom. Section 3.2 reports the residual diagnostics that motivate the regime structure and the Student- t errors, including the two-step fit of the degrees of freedom.

3.2 Volatility Regimes and Heavy-Tailed Shocks

For the heteroskedasticity identification I use $h = 4$ volatility regimes. Table 1 reports the dates. Each regime boundary is a boundary of one of the four balance-sheet episodes of Table A.1. The first break is the end of the QE purchases in 2014:M11. The second is the start of the COVID purchases in 2020:M3, and the third is the start of the QT2 runoff in 2022:M6. Each of the four regimes contains one balance-sheet episode, and none contains two.

¹²The ratio lies between zero and one and equals zero before the ON RRP facility became active. It enters in levels, with no log or inverse-hyperbolic transform.

¹³The hyperparameters are overall tightness 3, lag decay 0.5, sum-of-coefficients (μ) weight 1, and co-persistence (λ) weight 5. The dummy scales use each variable’s univariate $\text{AR}(p)$ residual standard deviation, and the prior precision stays invariant to units.

Table 1: Volatility regimes.

Regime	Start	End	Description
1	2008:M11	2014:M11	QE1 through QE3.
2	2014:M12	2020:M2	Between the QE purchases and the COVID purchases, and QT1.
3	2020:M3	2022:M5	COVID-19 QE.
4	2022:M6	2026:M6	QT2 and the months after it ended.

Note: End dates are inclusive. Each break is a date on the balance sheet, the end of the QE purchases, the start of the COVID purchases, and the start of the QT2 runoff. Each regime contains one balance-sheet episode of Table A.1. The QT2 runoff ended in 2025:M12, and the six months that follow are too few for a regime of their own.

The sample spans the financial crisis, COVID, and the post-2021 tightening. Figure 4 plots the reduced-form residuals u_t in (3) from estimating (1) with a single variance regime and Gaussian errors, without the identification restrictions, under the prior of Section 3.1. Each panel shows the posterior median residual of one equation in its own units. The four subsamples follow the dates of Table 1 and partition the residual sample for description only. The number above each subsample is its residual standard deviation.

The standard deviations vary across subsamples, and they do not move together across equations. The one-year yield is four times as volatile in the third subsample as in the first. The term spread barely moves across the four. Non-proportional patterns of this kind motivate the regime variances. Equation (4) identifies the columns of A_0 when no two structural shocks change their variances in the same proportion across regimes. At the posterior mode of the full model and across the retained draws, the estimated variance ratios are far from proportional.¹⁴

I measure each residual against the typical size of a shock in its own regime, and a few months still stand far outside the rest. The SOMA residual of April 2020 is 9.6 times that size, and the one-year yield residual of March 2020 is 5.5 times. Those are the months of the emergency purchases at the onset of COVID and of the emergency cuts to the zero lower bound.

Without the latent scales $v_{i,t}$ of (4) the estimated variance of a regime that holds one of these months reflects that month rather than the regime. The heteroskedasticity identification weakens with it. The latent scales give each shock in each month its own scale. The two months lose weight

¹⁴Table D.2 in Appendix D reports them.




Figure 4: Reduced-form residuals under a single variance regime, three equations.

Note. Each panel plots the posterior median of the reduced-form residual of one equation from the estimation of (1) with a single variance regime, Gaussian errors, no identification restrictions, and the prior of Section 3.1. Units are percent for SOMA and percentage points for the one-year yield and the term spread. Figure D.1 plots the other five equations. Dashed lines mark the regime breaks of Table 1. The shaded band in each subsample is plus and minus one subsample standard deviation of the series, and the number above each subsample is that standard deviation in the units of the panel. The residual sample runs from 2009:M5 to 2026:M6.

in the likelihood, and the regime variances keep their meaning. I fit the Student- t degrees of freedom by the two-step procedure of Brunnermeier et al. (2021) and use the fitted value of 5.84. Section D reports the diagnostics behind these numbers.

3.3 Identification

I identify the shocks with three sets of restrictions. Sign restrictions label six of the eight shocks. Zero restrictions hold at impact. The zeros bind under the average transmission, the responses evaluated at the average of the relative variances across regimes, and need not hold in the regime-specific responses (Appendix C.1). Narrative restrictions tie RFBS and OFBS to the episodes that generated them. All sign restrictions bind from impact through horizon five and must hold both under the average transmission and in every regime. I describe the two balance-sheet shocks first and the four remaining labeled shocks second. Table 2 collects the restrictions on the balance-sheet shocks, and Table 3 collects the restrictions on the demand, supply, monetary policy (MP), and EBP shocks.

The balance-sheet shocks. Table 2 collects the restrictions on RFBS and OFBS. The two shocks correspond to the two financing channels. I normalize the positive direction of both shocks as an expansion. A positive RFBS shock raises SOMA, raises $\log(\text{Reserves} + \text{ON RRP})$, and compresses the term spread, and the expansion settles in reserves. A positive OFBS shock moves the same

Table 2: Restrictions on the two balance-sheet shocks.

Response variable	RFBS	OFBS
<i>Panel A: sign and zero restrictions</i>		
One-year yield		
SOMA	+	+
log(Reserves + ON RRP)	+	+
ON RRP/(Reserves + ON RRP)	−	+
GDP	Z	Z
GDPDEF	Z	Z
Term spread	−	−
EBP		
<i>Panel B: narrative restrictions (window W, regime average)</i>		
Window W	2020:M4–2021:M2	2023:M9–2024:M7
Mean shock realization, eq. (B.1)	+	−
HD on log(R + RRP), eq. (B.2)	dominant	dominant

Note: A + or − is a sign restriction on the response of the row variable to the column shock, binding from impact through horizon five $[0, 5]$. Z is a zero at impact. Blank cells are unrestricted. Panel B ties each balance-sheet shock to its window. Over the window the mean realization of RFBS must be positive and the mean realization of OFBS must be negative, and the shock’s historical-decomposition contribution to log(Reserves + ON RRP) must exceed the sum of the contributions of all other labeled shocks.

three variables in the same direction, and the expansion settles in ON RRP. The ON RRP ratio separates the two columns. A reserve-financed balance-sheet expansion lowers the ratio (reserves raise the total while ON RRP is unchanged). An ON RRP-financed expansion raises the ratio directly. Without the ratio restriction the two columns would satisfy identical restrictions. The restrictions carry no episode label. QE and QT1 enter as positive and negative realizations of RFBS, and the QT2 drain enters as negative OFBS realizations.

One economic assumption generates the zero restrictions in Panel A of Table 2.

Assumption 1 (Balance-sheet shocks do not reach real activity and prices within the month). *The two balance-sheet shocks have no impact effect on real GDP or the GDP deflator.*

Assumption 1 follows the timing of [Leeper et al. \(1996\)](#), whose sluggish private sector responds to financial variables only with a delay. Production and final-goods pricing involve planning processes, and movements in interest rates and in the balance sheet reach output and prices smoothly over

time rather than within the month. [Christiano et al. \(2005\)](#) impose the same timing on the MP shock. Assumption 1 applies it to the two balance-sheet shocks instead, and the MP shock stays sign-restricted at impact. The balance sheet changes gradually after each program announcement ([Nakashima et al., 2024](#)), and its real effects plausibly arrive after the quantities accumulate. The timing argument binds more tightly for the quantity channel than for the rate channel.

Beyond impact, the GDP and deflator responses to both balance-sheet shocks stay unrestricted. The zeros bind on the average transmission, the response evaluated at the regime average of the relative variances. That average is the object I report. The regime-specific impact responses stay unrestricted ([Section 4](#)). The macroeconomic transmission is data-determined and can differ across the two channels.

I impose two narrative restrictions, one for each balance-sheet shock, and Panel B of [Table 2](#) summarizes them.¹⁵ Each adapts the single-date narrative sign restrictions of [Antolín-Díaz and Rubio-Ramírez \(2018\)](#) to a window average and ties the shock to the episode that generated it through two requirements. The mean realization of the shock over the window must carry the sign of its episode. The shock’s contribution to the unexpected movement of $\log(\text{Reserves} + \text{ON RRP})$ over the window must dominate the contributions of the other labeled shocks.¹⁶ The baseline measures the contribution by the multi-horizon interval historical decomposition of [Antolín-Díaz and Rubio-Ramírez \(2018\)](#). [Appendix B](#) states the two requirements formally.

I tie the RFBS shock to the first phase of the COVID purchases, April 2020 through February 2021, when the expansion settled predominantly in bank reserves. The window carries a positive sign. From March 2021 the inflows settled increasingly in MMFs and the ON RRP facility ([Section 2](#)). Including those months would risk misallocating the TGA-driven movement in $\log(\text{Reserves} + \text{ON RRP})$ to the RFBS shock.

¹⁵I also estimated the model with four narrative restrictions. The first addition anchors a negative mean RFBS realization over the QT1 window. The second anchors a positive mean OFBS realization over the ON RRP buildup, April 2021 through February 2022. The estimates stay almost identical, and the two restrictions below are enough in this sample.

¹⁶I impose the dominance on $\log(\text{Reserves} + \text{ON RRP})$ rather than on SOMA for two reasons. [Section 2](#) defines the two episodes by their effects on the liability side, and $\log(\text{Reserves} + \text{ON RRP})$ carries no zero restriction. The requirement discriminates the balance-sheet shock against every labeled shock, including the demand, supply, and MP shocks. On SOMA the zero restrictions would leave only the other balance-sheet shock as a competitor.

I tie the OFBS shock to the sharp ON RRP drain of September 2023 through July 2024. The window contains eleven months, the same length as the RFBS window. A positive OFBS shock is an expansion, the QT2 drain corresponds to negative OFBS realizations, and the window carries a negative sign. Over that window the ON RRP facility ran down steadily while reserves stayed roughly flat. This window starts more than a year after QT2 began in mid-2022, as Section 2 documents.¹⁷ Tying the restrictions to these phases, rather than to full episodes, keeps each inside the window where its channel was the active margin. Averaging over windows rather than single dates makes the restriction robust to the timing noise of individual months.

The demand, supply, MP, and EBP shocks. Table 3 collects the restrictions on the four remaining labeled shocks. I label them because the dominance requirement compares each balance-sheet shock against the other labeled shocks. The labels put demand, supply, MP, and the EBP shock among the competitors it must dominate over its window. Without the labels the comparison would run against the other balance-sheet shock alone. Demand, supply, and a conventional MP shock follow standard conventions. A demand shock raises the one-year yield, GDP, and the GDP deflator. A supply shock raises GDP and lowers the deflator. An MP shock raises the one-year yield and lowers GDP and the deflator. The EBP shock raises the EBP. Two economic assumptions generate the zero restrictions on these shocks.

Assumption 2 (Balance-sheet policy drives SOMA within the month). *Within the month the Fed moves its SOMA holdings only through its balance-sheet policy. The demand, supply, MP, and EBP shocks have no impact effect on SOMA.*

Assumption 2 sets four zeros on SOMA, and the reasons differ by shock. For the demand and

¹⁷Part of the ON RRP decline over the window coincides with the post-debt-ceiling rebuild of the Treasury General Account (TGA). The Treasury refilled the TGA through heavy bill issuance from June through October 2023, and the rebuild drew down ON RRP rather than reserves. This raises the question of whether to attribute the drain to QT2 or to the fiscal rebuild. The two are far apart in magnitude. From September 2023 onward the monthly decline in ON RRP is roughly three times the contemporaneous rise in the TGA. The rebuild does not account for the bulk of the drain, and among the six labeled shocks OFBS is the dominant driver of $\log(\text{Reserves} + \text{ON RRP})$ over the window. I read the restriction as identifying the dominant balance-sheet shock, while noting that the bill issuance associated with the rebuild plausibly amplified the same channel. Section 5.3 reports specifications that control the TGA and Treasury issuance directly.

Table 3: Restrictions on the demand, supply, MP, and EBP shocks.

Response variable	Demand	Supply	MP	EBP
One-year yield	+		+	
SOMA	Z	Z	Z	Z
log(Reserves + ON RRP)				
ON RRP/(Reserves + ON RRP)				
GDP	+	+	–	Z
GDPDEF	+	–	–	Z
Term spread				
EBP				+

Note: A + or – is a sign restriction on the response of the row variable to the column shock, binding from impact through horizon five [0, 5]. Z is a zero at impact. Blank cells are unrestricted. The regime covariances identify (A_0, Λ) and do not label the shocks, and the two unlabeled shocks carry no restriction.

supply shocks the zero follows the policy-equation exclusions of [Leeper et al. \(1996\)](#). In their system the policy instrument does not respond within the month to output and final-goods prices, the counterparts of GDP and the GDP deflator here. They give two grounds. The data for those series arrive only after the month ends, and a shift in the policy stance requires consultation and consensus building. The preannounced schedules and caps of asset purchases and runoff strengthen both grounds, and the within-month path of the balance sheet does not react to macroeconomic news.

For the EBP shock only the need for consultation and consensus building applies, since the corporate bond prices behind the EBP arrive within the month. The preannounced schedules and caps of purchases and runoff keep SOMA from responding to it within the month. The restriction constrains the impact responses of SOMA only, and GDP and the GDP deflator carry zeros under Assumptions 1 and 3 instead.

For the MP shock the zero reflects the separation of the two instruments. Since October 2008 the Fed has implemented the policy rate through administered rates on reserves and the ON RRP facility rather than through open market operations. A rate surprise requires no change in asset holdings within the month. An MP shock that moved SOMA contemporaneously would also conflate the two

instruments. The sample begins in November 2008 and lies entirely within this operating framework.

[Leeper et al. \(1996\)](#) find that movements in reserves and M1 in their sample arise mostly from policy accommodating shifts in private demand. An identification that treats reserves or M1 as a policy instrument mislabels those demand shifts as policy shocks. They call such identifications unreliable. The critique applies to any quantity that private demand can move. SOMA is not such a quantity. Only the Fed trades in its own portfolio, and under administered rates the accommodation of demand runs through the liability side rather than through asset holdings.

Assumption 3 (Shocks to the EBP do not reach real activity and prices within the month). *The EBP shock has no impact effect on real GDP or the GDP deflator.*

Assumption 3 is the identifying assumption of [Gilchrist and Zakrajšek \(2012\)](#), in whose quarterly VAR shocks to the EBP affect economic activity and inflation with a lag. Over a month the same assumption restricts less.

4 HARS-Z

This section presents the sampler that draws from the posterior of the heteroskedastic model of [Section 3](#) conditional on the sign, zero, and narrative restrictions of [Section 3.3](#). HARS-Z is the heteroskedastic absorb-and-rotate sampler with zero restrictions. I develop it by extending HARS, the heteroskedastic absorb-and-rotate sampler of [Kim and Zha \(2026\)](#), hereafter KZ) that combines shock heteroskedasticity with sign and narrative restrictions, to include zero restrictions. I adopt the notation of KZ throughout. Appendix C contains the full theoretical analysis. Readers focused on the estimates can note [Algorithm 1](#) and continue to [Section 5](#).

Throughout, $\theta = (A_0, \Lambda, A_+)$ collects the structural objects with $\Lambda = \{\Lambda_k\}_{k=1}^h$, $Q \in \mathcal{O}(n)$ denotes the rotation, and $SR(\theta, Q)$ is the indicator that the sign and narrative restrictions hold. HARS (KZ, Algorithm 1) combines likelihood-preserving posterior draws of the model parameters with rotations drawn from a well-defined probability distribution restricted by the sign and narrative restrictions, adopting a column-wise elliptical slice sampler (ESS; [Murray et al., 2010](#)). Sign and

narrative restrictions are inequalities. They truncate the conditional distribution of Q to a subset of positive Haar measure, and the posterior target in HARS (Appendix C.1) is well defined. A random rotation lands in the subset with positive probability (Figure 5, panel a). Appendix C.1 reviews the sampler and the repair steps in detail.




Figure 5: Why zero restrictions break HARS, and what HARS-Z does instead.

Note. Each sphere stands for $O(n)$. The shaded region is the set of rotations that satisfy the sign and narrative restrictions, and the dashed curve is the zero manifold $Q(\theta)$. In panel (a) HARS draws rotations from the starting point and keeps the ones that land in the region. In panel (b) a zero restriction adds the curve, a set of measure zero. A random rotation reaches it with probability zero, and every draw is discarded. In panel (c) HARS-Z builds each rotation on the curve and the region selects among them.

A zero restriction is not a truncation. A zero restriction on the impulse responses is a linear equality constraint on the j -th column q_j of Q . Following the notation of Arias et al. (2018, hereafter ARRW), the $z_j \times n$ matrix $Z_j(\theta)$ stacks the z_j zero restrictions on shock j as rows. The restrictions on shock j read $Z_j(\theta) q_j = 0$, the rows depend on the draw through A_0 and the moving-average coefficients, and the sampler must rebuild them at every iteration. The zeros bind under the average transmission only and need not hold in the regime-specific responses, where $\Lambda_k^{1/2}$ reweights the constrained column. Define the zero manifold

$$Q(\theta) = \{ Q \in O(n) : Z_j(\theta) q_j = 0 \text{ for every restricted } j \}. \quad (5)$$

Whenever some $z_j \geq 1$, $Q(\theta)$ is a lower-dimensional smooth manifold in $O(n)$ and $\mu(Q(\theta)) = 0$, where μ is Haar measure on $O(n)$, as ARRW note in motivating their Algorithm 2 (Figure 5, panel b).

One may be tempted to impose the zeros by adding them to $SR(\theta, Q)$ and running HARS

unchanged. The zeros lower the dimension of the admissible set $R(\theta) = \{Q \in O(n) : SR(\theta, Q) > 0\}$ instead of truncating it, and the condition $\mu(R(\theta)) > 0$ fails at every θ with at least one zero restriction. The restricted posterior target is not defined, the Gaussian coordinatization never reaches the zeros, and every step of KZ's Algorithm 1 fails with probability one, including its initialization. The failure originates in the measure rather than in the algorithm. Appendix C.2 shows it formally. The sampler must parameterize the zero-restricted set directly rather than reach it by sampling. It must build the rotation on $Q(\theta)$, and its slice-sampling updates must never leave it (Figure 5, panel c). The rest of this section constructs exactly such a parameterization.

HARS-Z replaces the QR map γ with a map $\gamma_z(\cdot; \theta)$, built on the null-space construction of ARRW's Algorithm 2, that sends a Gaussian matrix to a rotation on $Q(\theta)$ by construction. These are the built rotations of Figure 5, panel c. The state of the chain is the pair (θ, X) , where X is an $n \times n$ Gaussian matrix, the coordinates in which HARS runs its column-wise ESS update. The sampler recomputes the rotation $Q = \gamma_z(X; \theta)$ whenever it needs one. A structural update moves the manifold $Q(\theta)$, the recomputed rotation lies on the new manifold, and the zeros hold after every move. Appendix C.3 builds the map γ_z and establishes its properties.

HARS-Z targets the analogue of the posterior target in HARS on the zero-restricted set, with the uniform measure on $Q(\theta)$ in place of Haar measure on $O(n)$,

$$\pi_v(d\theta, dQ \mid y, SR) = \pi(\theta \mid y) \frac{\mathbf{1}\{Q \in R(\theta)\} \mu(dQ \mid Q(\theta))}{\mu(R(\theta) \mid Q(\theta))} d\theta, \quad (6)$$

where $\mu(\cdot \mid Q(\theta))$ denotes the uniform measure on $Q(\theta)$ and $R(\theta) = \{Q \in Q(\theta) : SR(\theta, Q) > 0\}$ now denotes the admissible set on the zero-restricted set. Generating the rotation through γ_z induces the null-space construction measure on $Q(\theta)$, not the uniform measure. Weighting each retained rotation by the volume element $v(Q; \theta)$ of the construction, normalized within each retained draw of θ , recovers the target π_v . The column-wise ESS update remains valid on $Q(\theta)$ and leaves the conditional target invariant. Appendix C.4 establishes these results.

The HARS-Z sampler, stated below as Algorithm 1, alternates a posterior draw of θ with a

rotation draw of X . The posterior draw is the sampler of Brunnermeier et al. (2021, hereafter BPSS), taken from KZ unchanged. Following KZ’s Theorem 1, partition the parameter collection as $z_0 = (A_0, \Lambda)$ and $z_+ = A_+$, and let s_0 collect everything held fixed in the block update. The first block moves by Metropolis–Hastings (MH) on the collapsed density with A_+ integrated out, and the companion block moves by a draw from its Gibbs conditional $p_+(A_+ | z_0, s_0)$. The rejection counts generated by the repair steps below are diagnostics of computational efficiency, not density terms (KZ, Proposition 5). Each repair returns a draw from the correct conditional distribution. The counts require no reweighting and enter no posterior computation. Appendices C.3 and C.4 state the lemma, propositions, and corollaries that Algorithm 1 cites.

Algorithm 1. (Heteroskedastic absorb-and-rotate sampler with zero restrictions (HARS-Z).)

Let the current admissible state be (θ, X) with $Q = \gamma_z(X; \theta)$ and $SR(\theta, Q) > 0$, and fix the tuning parameters $N_{\text{ess}}, N_\theta, N_Q$. For the initialization, construct an admissible rotation by a numerical search at the initial θ and convert it into Gaussian coordinates by Lemma 1(iii). One iteration proceeds as follows. Each step matches its counterpart in Algorithm 1 of KZ.

1. **Posterior draw with admissibility check.** Propose $z_0^{\text{prop}} = (A_0^{\text{prop}}, \Lambda^{\text{prop}})$ by the first-block MH step and set $z_0^\star = z_0^{\text{prop}}$ if accepted and $z_0^\star = z_0$ otherwise. Draw the companion block $z_+^\star \sim p_+(A_+ | z_0^\star, s_0)$ together with the latent Student- t scales, and form the completed draw $\theta^\star = (z_0^\star, z_+^\star)$. Rebuild the rotation at the new draw, $Q^\star = \gamma_z(X; \theta^\star)$. The rebuilt rotation satisfies every zero restriction by construction. Check the sign and narrative restrictions. If $SR(\theta^\star, Q^\star) > 0$, set $\theta \leftarrow \theta^\star$ and go to Step 3. Otherwise the rotation is stale. Go to Step 2a.
- 2a. **Posterior-draw repair by the joint procedure.** Repair on the completed draw, as required by KZ’s Theorem 1. For at most N_θ attempts, draw a completed candidate $\theta^{\text{cand}} = (z_0^{\text{cand}}, z_+^{\text{cand}})$ and retain it only if $SR(\theta^{\text{cand}}, \gamma_z(X; \theta^{\text{cand}})) > 0$. The first-block MH step, anchored at the last admissible state θ (not at the stale draw $\theta^\star$), proposes the candidate block z_0^{cand} , and $z_+^{\text{cand}} \sim p_+(A_+ | z_0^{\text{cand}}, s_0)$. On success, set $\theta \leftarrow \theta^{\text{cand}}$ and go to Step 3. If the N_θ attempts retain no candidate, discard all candidates and go to Step 2b, carrying the draw $\theta^\star$ that

triggered the repair. Failed candidates, and the stale draw $\theta^\star$ itself, are auxiliary proposals. The chain never retains them.

- 2b. **Global rotation repair.** Hold $\theta^\star$ fixed. For $j = 1, \dots, N_Q$, draw a fresh Gaussian matrix $G^{(j)}$ and propose $Q^{(j)} = \gamma_z(G^{(j)}; \theta^\star)$, the analogue on $Q(\theta^\star)$ of the independent Haar draws in Step 2b of KZ’s Algorithm 1. Let $j^\star$ be the first index with $SR(\theta^\star, Q^{(j^\star)}) > 0$. If such an index exists, set $(\theta, X) \leftarrow (\theta^\star, G^{(j^\star)})$ and go to Step 3. If no index succeeds after N_Q attempts, declare sampler failure and proceed to Step 5.
3. **Rotation draw.** From the admissible state, run N_{ess} column-wise ESS updates of X at fixed θ (Proposition 2), producing Gaussian states $X^{(1)}, \dots, X^{(N_{\text{ess}})}$ and rotations $Q^{(r)} = \gamma_z(X^{(r)}; \theta)$ with $SR(\theta, Q^{(r)}) > 0$ for every r . Retain all N_{ess} rotations together with θ for posterior impulse-response computation.
4. **Advance the state.** Set the next state to $(\theta, X^{(N_{\text{ess}})})$ and return to Step 1.
5. **Sampler failure.** If a posterior move produces a draw $\theta^\star$ for which neither the joint posterior-draw repair nor the global rotation repair restores sign admissibility, the sampler cannot advance. Increase N_θ or N_Q , improve the repair search, or revisit the restriction set.

Retaining N_{ess} rotations per structural draw divides the cost of the structural update across several rotation draws. The ESS updates evaluate only the admissibility indicator, not the likelihood. Appendix C.4 establishes the validity of Algorithm 1 step by step. Posterior inference weights each retained rotation by its volume element, normalized within each retained draw of θ , and reports weighted medians and bands. Appendix C.5 states the estimator. HARS-Z nests HARS. With an empty zero set, γ_z reduces to the QR map γ , $Q(\theta) = O(n)$, and the construction measure $\mu_z(\cdot; \theta)$ is Haar measure. The volume weight is identically one, the weighted summaries reduce to the unweighted ones, and the sampler reduces to KZ’s Algorithm 1.

5 Results

I estimate the baseline posterior with HARS-Z. Appendix D reports the sampler settings and the diagnostics.

5.1 Baseline Impulse Responses

Figures 6 and 7 report the impulse responses to the two balance-sheet shocks under the baseline identification of Section 3.3, with the relative variances set to their average of one across regimes. The sign and narrative restrictions of Section 3.3 define each shock in its expansionary direction, and the identification is unchanged. Each figure shows one shock in its contractionary direction, the direction of the QT episodes. To make the two shocks directly comparable, I normalize each shock to a one-percent impact decline in SOMA holdings. The decline continues after impact. Figure E.1 in Appendix E reports the responses to the demand, supply, MP, and EBP shocks.

Figure 6 plots the responses to the RFBS shock. The shaded regions are the 68% credible bands. The restricted responses hold by construction. Panel A of Table 2 restricts the expansionary direction. Under the contractionary display, SOMA, $\log(\text{Reserves} + \text{ON RRP})$, the ON RRP ratio, and the term spread move with the signs of Panel A reversed. The zero restrictions hold GDP and the GDP deflator at zero on impact. The economic content lies in the GDP and GDP-deflator responses beyond impact and in the EBP response, where no restriction binds.

A negative RFBS shock lowers GDP persistently. The GDP response turns negative at five months and falls through the horizon, and the deflator response turns negative at two months. The 68% bands exclude zero from twelve months onward for GDP and from sixteen months onward for the deflator. The same shock raises the EBP. The response rises over the first fifteen months and stays positive through the horizon, and the 68% bands exclude zero from two months onward.

Figure 7 plots the responses to the OFBS shock. The restricted responses again hold by construction and reverse the signs of the OFBS column in Panel A of Table 2. The economic content again lies in the GDP and GDP-deflator responses beyond impact and in the EBP response, where




Figure 6: Impulse responses to the RFBS shock.

Note. Impulse responses to the RFBS shock, normalized to a one-percent impact decline in SOMA holdings. I evaluate the responses at the regime average of the relative variances. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. Vertical axes are in basis points for the one-year yield and the term spread, percent for SOMA, log(Reserves + ON RRP), GDP, and GDPDEF, and percentage points for the ON RRP ratio and the EBP.

no restriction binds.

A negative OFBS shock lowers the deflator. The response turns negative at one month and falls through the horizon, and the 68% bands exclude zero from six months onward. The GDP bands include zero after the first month, and the EBP bands include zero at every horizon.

Interpreting the two shocks. Neither figure describes a single episode. The linear model has one response function for each shock, and realizations of both signs enter its estimation. QE and QT1 enter the RFBS shock as positive and negative realizations, and the ON RRP buildup and the QT2




Figure 7: Impulse responses to the OFBS shock.

Note. Impulse responses to the OFBS shock, normalized to a one-percent impact decline in SOMA holdings. The figure shows the shock in the direction of the QT2 drain. I evaluate the responses at the regime average of the relative variances. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. Vertical axes are in basis points for the one-year yield and the term spread, percent for SOMA, $\log(\text{Reserves} + \text{ON RRP})$, GDP, and GDPDEF, and percentage points for the ON RRP ratio and the EBP.

drain enter the OFBS shock the same way (Section 3.3). Realizations of both signs also fall outside those episodes. The contraction display is a sign convention. I read each plotted response as the average over both directions and not as the response to the QT episode alone.

5.2 What Carries the Two Responses

This subsection asks what carries the effects of each balance-sheet shock. For the RFBS shock I hold the EBP response at zero. For the OFBS shock I add interest on reserves less the one-month

bill rate to the system.

5.2.1 The Reserve-Financed Contraction and the EBP

I hold the EBP response to the RFBS shock at zero and compare the responses of GDP and the deflator with and without it. Following BPSS, who keep the federal funds rate constant with a sequence of monetary policy shocks, I keep the EBP response at zero with a sequence of EBP shocks. I choose the EBP shocks draw by draw so that the EBP response is zero in every month from 0 to 60. In 174 of the 10,000 draws the shocks this requires grow without bound. I drop those draws and renormalize the weights of the other 9,826.¹⁸ Under Assumption 3 the EBP shock does not move GDP or the deflator on impact, and their counterfactual responses equal the actual ones on impact.




Figure 8: Responses to the RFBS shock with the EBP held at zero.

Note. Responses to the RFBS shock, normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid blue lines and blue regions show the RFBS shock alone. Dashed green lines and green regions show the RFBS shock combined with the EBP shocks that keep the EBP response at zero in every month. Lines are posterior medians and shaded regions are 68% credible bands over the 9,826 draws in which those shocks stay finite. Vertical axes are in percent for GDP and GDPDEF and in percentage points for the EBP.

With the EBP response held at zero, the declines in GDP and the deflator after a reserve-financed contraction are no longer present (Figure 8). Neither band lies below zero at any horizon. A

¹⁸In each posterior draw, let $r_x(h)$ be the response of variable x to the RFBS shock at month h . Let $b_x(h)$ be its response to an EBP shock of one standard deviation. Both responses use the average relative variances. Adding EBP shocks $\zeta(0), \dots, \zeta(60)$ to the RFBS shock gives the counterfactual response $c_x(h) = r_x(h) + \sum_{j=0}^h b_x(h-j) \zeta(j)$. I choose $\zeta(h)$ so that $c_{EBP}(h) = 0$, one month at a time from $h = 0$. The penalized offset of Figure F.2 instead chooses $\zeta(0), \dots, \zeta(60)$ to minimize $\sum_{h=0}^{60} c_{EBP}(h)^2 + \rho \sum_{h=0}^{60} \zeta(h)^2$, with $\rho = 10^{-3} \max_h r_{EBP}(h)^2$, and keeps every draw.

penalized offset that keeps all 10,000 draws gives the same result (Figure F.2 in Appendix F). Figure F.1 reports the other five variables.

The counterfactual attributes the effects of a reserve-financed contraction on GDP and the price level to the rise in the EBP, which lowers economic activity (Gilchrist and Zakrajšek, 2012). Section 2.3 offers two channels for that rise. A reserve-financed contraction would lower bank income and with it bank capital (link 3 of Figure 3), and it would tighten the leverage constraint of banks (link 4). Either channel would raise the risk premium banks require to bear credit risk. The estimates are consistent with both channels and do not separate them.

5.2.2 The ON RRP-Financed Contraction and the Bill Rate

MMFs move the ON RRP balance to whichever of the facility and bills pays more (Section 2.3). If bill rates move against the rates the Fed pays, MMFs move the balance, and the OFBS shock may pick up that movement rather than the financing itself. I add interest on reserves (IOR) less the one-month bill rate to the system as a ninth variable, with no restriction on it.¹⁹

The reserve-financed contraction keeps its effects on GDP, the deflator, and the EBP (Figure 9). The ON RRP-financed contraction loses its effect on the deflator, whose band includes zero at every horizon. With the spread in the system, the VAR attributes to the spread the movements in the ON RRP balance that follow it, and the OFBS shock carries what remains. The baseline cannot separate the deflator decline after an ON RRP-financed contraction from MMFs moving the balance as the spread moves.

5.3 Robustness

Eleven specifications each change one element of the baseline of Section 5.1 and keep every other element. Three remove zero restrictions, Assumption 1, Assumption 2, and both. Three change the variance structure. They use one regime variance (homoskedasticity) in place of four, and Gaussian

¹⁹I use IOR in place of the ON RRP rate because the ON RRP rate starts in September 2013. A spread on the ON RRP rate would cut the sample at October 2013 and drop most of the first variance regime, over which reserves financed the balance sheet. IOR and the one-month bill rate cover the whole sample.




Figure 9: Responses to the two balance-sheet shocks with the IOR less one-month bill spread.

Note. Rows report the RFBS and OFBS shocks, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines and shaded regions in the color of each shock are the baseline posterior medians and 68% credible bands of Figures 6 and 7. Dashed green lines and green regions are the posterior medians and 68% credible bands in the nine-variable system. Vertical axes are in percent for GDP and GDPDEF and in percentage points for the EBP.

errors in place of the Student- t with latent scales. The third sets the regime breaks at the exits from and the return to the zero lower bound (2015:M11, 2020:M3, and 2022:M2) in place of the balance-sheet episodes. Five change the regressors. Two add exogenous controls, the TGA ratio in one and the bill share and coupon issuance in the other. One puts industrial production and the CPI in place of the interpolated GDP and deflator. One puts the two-year Treasury yield in place of the one-year yield. One puts the corporate bond spread of [Gilchrist and Zakrajšek \(2012\)](#), from which the EBP is built, in place of the EBP. Table 4 gives, for each specification, the month from which the 68% bands exclude zero for the responses of GDP, the deflator, and the EBP to each of the two shocks. Appendix G plots those responses.

The main results survive every change except one. In all eleven specifications a reserve-financed contraction lowers GDP and raises the credit spread, and an ON RRP-financed contraction lowers the deflator and does not move the credit spread. In ten of the eleven the ON RRP-financed contraction

Table 4: Robustness of the GDP, deflator, and EBP responses.

Specification	RFBS			OFBS		
	GDP	Deflator	EBP	GDP	Deflator	EBP
Baseline	11	19	2	none	6	none
Without Assumption 1	33	none	3	none	13	none
Without Assumption 2	11	16	2	none	6	none
Without any zero	34	none	3 to 46	none	12	none
One variance regime	35	24	8	none	9	none
Gaussian errors	23	27	7	none	6	none
Regime breaks at the lower bound	15	25	6	none	7	none
TGA control	6	28	2	none	8	none
Issuance controls	9	12	2 to 29	none	9 to 28	none
Industrial production and CPI	17	none	3 to 54	none	5	none
Two-year yield as the policy rate	13	15	2	13 to 52	8	none
GZ credit spread in place of the EBP	21	21	13 to 31	none	6	none

Note: Each cell is the first month, counting from impact as month zero, at which the 68% band excludes zero with the sign the baseline response has. That sign is a decline for GDP and the deflator and a rise for the EBP. The band excludes zero through month sixty unless a second month closes the range. “none” means the band excludes zero with that sign at no horizon. Under the industrial-production specification the GDP and deflator columns read industrial production and the CPI, and under the GZ specification the EBP columns read the GZ credit spread. Without any zero the OFBS EBP band excludes zero at month sixty alone. In the baseline and in five specifications the OFBS GDP band excludes zero with a positive median over the first one to two months and at no later horizon.

does not move GDP either. With the two-year yield as the policy rate it lowers GDP as well. The two-year yield carries more of the expected policy path than the one-year yield, and the OFBS shock absorbs part of the 2022 to 2024 rate tightening that coincided with the drain. The specifications otherwise differ only in when each response appears, and [Table 4](#) gives the month.

Two things change. The reserve-financed decline in the deflator is the fragile result. It disappears when the impact zeros on GDP and the deflator are dropped and when the CPI replaces the deflator. Without the impact zeros the RFBS shock also carries a positive GDP response at impact, and the decline appears in the third year. I read that positive response as the state of the economy around the balance-sheet actions, the COVID months in the main, rather than as transmission. The impact zeros are what separates the two. The other change is smaller. With Treasury issuance as a control, the RFBS rise in the EBP and the OFBS decline in the deflator last about two years instead of the

full horizon. Bill issuance is the route the ON RRP drain takes to the Treasury ([Section 2](#)), and a control that absorbs the route absorbs part of the OFBS shock.

5.4 Variance and Historical Decompositions

The impulse responses show how the effects of a balance-sheet policy of a given size differ with its financing. They do not show how much of the variation in GDP and the deflator the two balance-sheet shocks account for. This subsection reports two decompositions that answer that question. The forecast-error variance decomposition gives the share of the forecast variance at a given horizon. The historical decomposition gives the contribution of each shock to the unexpected change in a variable over an episode. Both use the baseline draws of [Section 5.1](#) and weight the draws as the impulse responses do.


Forecast-error variance decomposition. [Figure 10](#) reports the median share of each labeled shock at every horizon from 0 to 60 months. The shares of both balance-sheet shocks in the forecast variance of GDP and the deflator rise with the horizon. Both shares start at zero and stay below 3% over the first year. They rise steadily after it. At five years the RFBS shock accounts for 9% of GDP and 8% of the deflator, and the OFBS shock for 5% of GDP and 9% of the deflator. The RFBS share of the EBP rises with the horizon as well. From one year on it is the largest share after that of the EBP shock itself.

Historical decomposition. The historical decomposition splits the accumulated forecast error of a variable over a window among the shocks that the sample realizes inside it. Let $W = \{t_0 + 1, \dots, t_1\}$ and let $\mathcal{I}_{t_0}$ collect the observables through t_0 . The contribution of labeled shock j to variable i is

$$\text{HD}_{j \rightarrow i}(W) = \sum_{\ell=0}^{t_1-t_0-1} [\Theta_{s_{t_1-\ell}, \ell}(Q)]_{ji} \tilde{\varepsilon}_{j, t_1-\ell}, \quad \sum_{j=1}^n \text{HD}_{j \rightarrow i}(W) = y_{i, t_1} - \mathbb{E}[y_{i, t_1} \mid \mathcal{I}_{t_0}], \quad (7)$$

where the regime index follows the date of the shock. [Figure 11](#) reports (7) over the four volatility regimes, as posterior medians with 68% intervals, for the six labeled shocks. I split the last regime

Figure 10: Forecast-error variance decomposition.



Note. Posterior median share of each labeled shock in the forecast-error variance of the variable, in percent, over the retained draws. The shares are under the average transmission, with the relative variances at their average of one across regimes. Horizons are in months. Table H.1 reports the shares at 36 months with their 68% intervals.

at 2023:M5, the month before the ON RRP balances began to fall. Figure I.1 in Appendix I reports the same decomposition for the other five variables.

Reserves mainly financed the balance sheet over the first two regimes (Figure 1), and the RFBS shock is the largest contributor to the GDP, deflator, and EBP forecast errors in both. Over Regime 1, the QE years, it pushes GDP and the deflator above their forecasts and the EBP below. Over Regime 2, which holds QT1, it pulls GDP and the deflator below their forecasts and pushes the EBP above. All six of these bands exclude zero. The OFBS bands include zero in both regimes.

Over Regime 3 the ON RRP balances build up to near \$2 trillion (Figure 1) while the COVID purchases add reserves as well, and neither liability mainly financed the purchases. The supply shock is the largest contributor to the GDP and deflator forecast errors, and the EBP shock to the EBP forecast error. The RFBS shock pushes GDP and the deflator above their forecasts, and both bands exclude zero. Its EBP band includes zero. The OFBS shock pushes the deflator above its forecast, and the band excludes zero. Its GDP band includes zero.

The split separates the first year of QT2, in which the balances barely move, from the months in which they fall. Over Regime 4 as a whole, no labeled shock has a band that excludes zero for GDP or the deflator. Over the first year the MP shock is the largest contributor to the GDP and deflator forecast errors, and neither band excludes zero. Over the drain the OFBS shock is the largest contributor to the GDP and deflator forecast errors. It pulls both below their forecasts, and both




Figure 11: Historical decomposition by volatility regime.

Note. Contribution of each labeled shock to the accumulated forecast error of the variable over the window. Units are percent for GDP and the deflator and percentage points for the EBP. Bars are posterior medians and whiskers are 68% credible intervals, both over the retained draws. The first four windows are the volatility regimes of Table 1, and the last two split the fourth at 2023:M5. Regime 1 opens at 2009:M5, the first estimable month under the six estimation lags. The two unlabeled shocks account for the remainder.

bands exclude zero. Its EBP band includes zero. The RFBS shock contributes nothing to any of the three variables over the drain.

Interpreting the decompositions. Each balance-sheet shock is the main source of the unexpected movements in GDP and the deflator when its own liability mainly financed the balance sheet. Reserves mainly financed QE and QT1, and the RFBS shock moved GDP, the deflator, and the EBP in both. GDP and the deflator rose when the EBP fell and fell when it rose, the pattern the counterfactual of [Section 5.2.1](#) attributes to the EBP. ON RRP mainly financed the drain, and the

OFBS shock moved GDP and the deflator without moving the EBP.

Over the drain MMFs also moved their balances from ON RRP to bills, and [Section 5.2.2](#) shows that the OFBS effect on the deflator does not survive a control for that reallocation. In the nine-variable system of [Section 5.2.2](#), the RFBS shock keeps its sign and remains the largest contributor to the GDP and deflator forecast errors over QE and QT1. The OFBS contribution to the deflator over the drain falls to near zero.

The shares in [Figure 10](#) say how much the financing matters on average. The two balance-sheet shocks account for little of the forecast variance of GDP and the deflator in the first year. Their shares grow with the horizon to about a tenth at five years. The RFBS share of the EBP grows faster and reaches 18% within two years, close to the share of the EBP shock itself. The balance-sheet shocks move the risk premium more than output and prices, and they move it through reserve financing.

6 Conclusion

Since 2008 the Fed has financed its balance-sheet operations with two overnight liabilities. Reserves financed the purchases after the financial crisis and the runoff of QT1. Both liabilities financed the COVID purchases, and ON RRP financed most of the runoff of QT2, which left reserves nearly intact ([Figure 1](#)). A reserve-financed operation changes bank income and, through it, bank capital. An ON RRP-financed operation changes what MMFs hold, not what they earn. This paper asks whether the transmission of balance-sheet policy to real GDP and the price level depends on which liability finances it.

Both balance-sheet contractions lower the deflator, and only the reserve-financed contraction lowers GDP. The EBP says why. A reserve-financed operation moves the EBP, a contraction raising it and an expansion lowering it, and an ON RRP-financed operation does not. Banks with more capital require a smaller risk premium to bear credit risk, and that premium moves the cost of credit for firms. At three years the two balance-sheet shocks account for a tenth of the forecast variance of GDP and the deflator. They account for a fifth of the forecast variance of the EBP, mostly through

the RFBS shock. In the sample the RFBS shock is the largest contributor to the GDP, deflator, and EBP forecast errors over the two regimes that reserves financed. The OFBS shock is the largest contributor to the deflator forecast error over the ON RRP drain of QT2.

The empirical design combines heteroskedasticity, sign restrictions, narrative restrictions, and zero restrictions. The heteroskedastic likelihood separates the model shocks, and the sign and narrative restrictions label the two balance-sheet shocks. HARS-Z imposes the impact-zero restrictions on that combination, and the rotation satisfies every zero restriction at every draw.

Which liability finances an operation decides whether the operation reaches output. Both financings move the price level, and only the reserve-financed one moves output. The Fed sets the terms that govern the split between reserves and ON RRP. Within those terms the composition of Treasury issuance, bank regulation, and the portfolio choices of MMFs affected which liability each episode filled or drained ([Section 2](#)). The asset side of the balance sheet moved alike in the two expansions and in the two contractions, and the liability side did not. The size of the balance sheet does not summarize what a balance-sheet policy does.

References

- Adrian, T., Erceg, C. J., Kolasa, M., Linde, J. L., and Zabczyk, P. (2025). Macroeconomic and fiscal consequences of quantitative easing.
- Afonso, G., Cipriani, M., and La Spada, G. (2025). Bank regulation and monetary policy: The role of non-bank financial institutions. Staff Reports 1041, Federal Reserve Bank of New York. First issued December 2022, revised December 2025.
- Antolín-Díaz, J. and Rubio-Ramírez, J. F. (2018). Narrative sign restrictions for SVARs. *American Economic Review*, 108(10):2802–29.
- Arias, J. E., Rubio-Ramírez, J. F., and Shin, M. (2026). Are fiscal transfers inflationary? Working paper, Federal Reserve Bank of Philadelphia.
- Arias, J. E., Rubio-Ramírez, J. F., and Waggoner, D. F. (2018). Inference based on structural vector autoregressions identified with sign and zero restrictions: Theory and applications. *Econometrica*, 86(2):685–720.
- Becker, B. and Ivashina, V. (2014). Cyclicity of credit supply: Firm level evidence. *Journal of Monetary Economics*, 62:76–93.
- Board of Governors of the Federal Reserve System (2026). Financial accounts of the united states: Flow of funds, balance sheets, and integrated macroeconomic accounts, second quarter 2026. Statistical Release Z.1, <https://www.federalreserve.gov/releases/z1/>. Release of September 11, 2026; tables F.521 and S.123.
- Brunnermeier, M., Palia, D., Sastry, K. A., and Sims, C. A. (2021). Feedbacks: Financial markets and economic activity. *American Economic Review*, 111(6):1845–79.
- Chen, H., Cúrdia, V., and Ferrero, A. (2012). The macroeconomic effects of large-scale asset purchase programmes. *The Economic Journal*, 122(564):F289–F315.

- Christiano, L. J., Eichenbaum, M., and Evans, C. L. (2005). Nominal rigidities and the dynamic effects of a shock to monetary policy. *Journal of Political Economy*, 113(1):1–45.
- Copeland, A., Duffie, D., and Yang, Y. D. (2021). Reserves were not so ample after all. Staff Reports 974, Federal Reserve Bank of New York.
- Crawley, E., Gagnon, E., Hebden, J., and Trevino, J. (2022). Substitutability between balance sheet reductions and policy rate hikes: Some illustrations and a discussion. FEDS notes, Board of Governors of the Federal Reserve System.
- d’Avernas, A., Hubert de Fraisse, A., Ning, L., and Vandeweyer, Q. (2024). The fiscal cost of quantitative easing. Working paper, Social Science Research Network.
- Diamond, W. F., Jiang, Z., and Ma, Y. (2024). The reserve supply channel of unconventional monetary policy. *Journal of Financial Economics*, 159:103581.
- Federal Open Market Committee (2019). Statement regarding monetary policy implementation. Board of Governors of the Federal Reserve System, October 11. Available at <https://www.federalreserve.gov/newsevents/pressreleases/monetary20191011a.htm>.
- Federal Open Market Committee (2022a). Plans for reducing the size of the federal reserve’s balance sheet. Board of Governors of the Federal Reserve System, May 4. Available at <https://www.federalreserve.gov/newsevents/pressreleases/monetary20220504b.htm>.
- Federal Open Market Committee (2022b). Principles for reducing the size of the federal reserve’s balance sheet. Board of Governors of the Federal Reserve System, January 26. Available at <https://www.federalreserve.gov/newsevents/pressreleases/monetary20220126c.htm>.
- Federal Open Market Committee (2024). Statement regarding reinvestment of principal payments from treasury securities, agency debt, and agency mortgage-backed securities. Federal Reserve Bank of New York, Operating Policy, May 1. Available at https://www.newyorkfed.org/markets/opolicy/operating_policy_240501.

- Federal Reserve Bank of New York (2026). Primary dealers list. <https://resources.newyorkfed.org/newsevents/news/markets/2026/an260115p>. Announcement of January 15, 2026.
- Fricke, D., Greppmair, S., and Paludkiewicz, K. (2026). Excess reserves and monetary policy tightening. *Review of Finance*, 30(5):1777–1816.
- Gertler, M. and Karadi, P. (2015). Monetary policy surprises, credit costs, and economic activity. *American Economic Journal: Macroeconomics*, 7(1):44–76.
- Gilchrist, S. and Zakrajšek, E. (2012). Credit spreads and business cycle fluctuations. *American Economic Review*, 102(4):1692–1720.
- Glover, A. and Oliyide, J. (2024). Why haven’t recent rate increases slowed the economy more? Look to unusually low private lending spreads. Economic bulletin, Federal Reserve Bank of Kansas City.
- He, Z. and Krishnamurthy, A. (2013). Intermediary asset pricing. *American Economic Review*, 103(2):732–770.
- Kim, D. and Zha, T. (2026). Sharpening economic interpretation with HARS. Working Paper 35483, National Bureau of Economic Research.
- Krishnamurthy, A. and Vissing-Jorgensen, A. (2012). The aggregate demand for treasury debt. *Journal of Political Economy*, 120(2):233–267.
- Kumhof, M. and Salgado-Moreno, M. (2024). Quantitative easing and quantitative tightening: The money channel. Staff Working Paper 1090, Bank of England.
- Lanne, M., Lütkepohl, H., and Maciejowska, K. (2010). Structural vector autoregressions with markov switching. *Journal of Economic Dynamics and Control*, 34(2):121–131.
- Leeper, E. M., Sims, C. A., and Zha, T. (1996). What does monetary policy do? *Brookings Papers on Economic Activity*, 1996(2):1–78.

- Lloyd, S. and Ostry, D. (2024). The asymmetric effects of quantitative tightening and easing on financial markets. *Economics Letters*, 238:111722.
- Logan, L. K. (2024). Opening remarks at panel on market monitoring and the implementation of monetary policy. Annual meetings of the International Banking, Economics and Finance Association and the American Economic Association, San Antonio, January 6. Available at <https://www.dallasfed.org/news/speeches/logan/2024/lk1240106>.
- Lu, Y. and Valcarcel, V. J. (2024). A tale of two tightenings. *Journal of Economic Dynamics and Control*, 166:104906.
- Murray, I., Adams, R. P., and MacKay, D. J. C. (2010). Elliptical slice sampling. In *Proceedings of the 13th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 9 of *JMLR Workshop and Conference Proceedings*, pages 541–548.
- Nakashima, K., Shibamoto, M., and Takahashi, K. (2024). Identifying quantitative and qualitative monetary policy shocks. *International Journal of Central Banking*, 20(3):1–61.
- Rigobon, R. (2003). Identification through heteroskedasticity. *Review of Economics and Statistics*, 85(4):777 – 792. Cited by: 459.
- Rubio-Ramírez, J. F., Waggoner, D. F., and Zha, T. (2010). Structural vector autoregressions: Theory of identification and algorithms for inference. *The Review of Economic Studies*, 77(2):665–696.
- Sims, C. A. and Zha, T. (1998). Bayesian methods for dynamic multivariate models. *International Economic Review*, 39(4):949–968.
- Sims, E. and Wu, J. C. (2021). Evaluating central banks’ tool kit: Past, present, and future. *Journal of Monetary Economics*, 118:135–160.
- Smith, A. L. and Valcarcel, V. J. (2023). The financial market effects of unwinding the federal reserve’s balance sheet. *Journal of Economic Dynamics and Control*, 146:104582.

- Swanson, E. T. (2021). Measuring the effects of federal reserve forward guidance and asset purchases on financial markets. *Journal of Monetary Economics*, 118:32–53.
- U.S. Securities and Exchange Commission (2026). Rule 2a-7, money market funds, 17 CFR 270.2a-7. Code of Federal Regulations. Paragraph (a)(14), definition of a government money market fund.
- Waller, C. J. (2024). Thoughts on quantitative tightening. Speech at the 2024 U.S. Monetary Policy Forum, New York, March 1. Available at <https://www.federalreserve.gov/newsevents/speech/waller20240301a.htm>.
- Wei, B. (2022). Quantifying “quantitative tightening” (QT): How many rate hikes is QT equivalent to? Working Paper 2022-8, Federal Reserve Bank of Atlanta.
- Wong, R. (2021). The borrower of last resort: What explains the rise of ON RRP facility usage? Economic Brief 21-43, Federal Reserve Bank of Richmond.

A Balance-Sheet Data

This appendix collects the episode-level balance-sheet evidence discussed in Section 2. Table A.1 reports the Fed’s flows by episode. Table A.2 reports the MMF and hedge fund balance sheets over the same windows, and Figure A.1 reports who held the ON RRP balance.

Table A.1: Fed balance-sheet flows by episode, USD billions.

	QE	QT1	COVID QE	QT2
	2008-12 → 2014-11	2017-10 → 2019-08	2020-03 → 2022-03	2022-06 → 2025-12
<i>Asset side</i>				
Total assets	+2,266	−690	+4,315	−2,352
MBS	+1,725	−267	+1,340	−666
SOMA Treasury	+1,985	−379	+3,095	−1,559
coupons	+2,004	−381	+3,087	−1,447
T-bills	−18	+2	+8	−112
<i>Liability side</i>				
Total liabilities	+2,266	−690	+4,315	−2,352
Reserves	+1,765	−709	+1,998	−312
ON RRP	+107	−115	+1,588	−2,152
TGA	+12	−43	+242	+124
Other	+381	+177	+487	−11
<i>Absorption ratios (% of total asset change)</i>				
Reserves	77.9	102.8	46.3	13.3
ON RRP	4.7	16.7	36.8	91.5

Notes: Each value is the monthly average in the last month of the window less the monthly average in the first month. The asset side uses the Federal Reserve H.4.1 series WALCL, WSHOMCB, TREAST, and WSHOBL. Coupons are the Treasury securities other than bills, including TIPS and FRNs. The liability side uses the H.4.1 series WRESBAL and WTREGEN and the daily ON RRP series RRPONTSYD. “Other” is total assets less reserves, ON RRP, and the TGA, and it holds currency in circulation and residual Fed liabilities (capital, foreign and other deposits). Total liabilities equal total assets by accounting identity.

Table A.2 records the ON RRP route. When MMFs take cash out of the facility, the offsetting item sits on their own asset side, in Treasury securities and in repo with private borrowers. Total assets need not fall. The hedge fund rows record who borrowed the repo cash and what they held. The issuance rows record the coupons and bills that the Treasury sold over each window, which is the supply MMFs and hedge funds absorbed.

Figure A.1 records who held the ON RRP balance. The MMF categories are the MMFs’ own

Table A.2: Money market fund and hedge fund balance sheets by episode, USD billions.

	QE 2008-12 → 2014-11	QT1 2017-10 → 2019-08	COVID QE 2020-03 → 2022-03	QT2 2022-06 → 2025-12
<i>ON RRP facility</i>				
All counterparties	+149	−304	+1,870	−1,873
of which MMFs		−286	+1,656	−1,681
<i>Money market fund assets</i>				
Treasury securities		+222	+719	+1,979
Private repo		+609	−605	+2,251
of which through FICC		+173	−106	+1,210
Total investments		+803	+1,058	+3,316
<i>Hedge funds</i>				
Government debt exposure		+953	−493	+2,497
Repo borrowing		+610	−212	+2,289
<i>Treasury issuance</i>				
Net bill issuance	−564	+530	+1,364	+2,874
Net coupon issuance	+7,163	+1,419	+5,006	+4,092

Notes: Changes over the episode windows of Table A.1. Money market fund, facility, and issuance rows are month-end levels at the last month of the window less the level at the last month-end before it. The base of the QT1 window, 2017:M9, is a quarter end, and the facility carries its quarter-end spike on that date, which is why the QT1 facility cell is larger than the change in monthly means that Table A.1 reports. Form N-MFP is a month-end snapshot. Hedge fund rows are quarter-ends that bracket the window. Blank cells have no observation at the base date. Form N-MFP begins in 2010:M11 and Form PF in 2013, both after the QE window opens. MMFs are money market funds. Private repo is their repo with every counterparty other than the Fed. Total investments are their portfolio holdings and exclude cash and other assets. FICC is the Fixed Income Clearing Corporation, where hedge funds borrow through sponsored repo. Government debt exposure is hedge funds' gross exposure to U.S. government debt, and repo borrowing is the cash they borrow in repo, both from Form PF. Treasury securities include coupons with remaining maturity under thirteen months. Sources: Federal Reserve daily ON RRP series (RRPONTSYD) and SEC Form N-MFP through the OFR Money Market Fund Monitor. Hedge fund rows come from SEC Form PF through the OFR Hedge Fund Monitor, and issuance rows from the Monthly Statement of the Public Debt.

Form N-MFP categories, pooled as the Office of Financial Research pools them, with Treasury, Government/Agency, and Exempt Government funds as government funds. The non-MMF remainder is the amount the New York Fed reports for government-sponsored enterprises, banks, and primary dealers. The series starts in 2016:M1. The Fed ran term reverse repurchase operations at quarter ends in 2014 and 2015, and MMFs report those alongside the overnight facility while the overnight balance excludes them.




Figure A.1: The ON RRP balance by holder.

Note. Top panel: the month-end ON RRP balance in USD billions, split into government money market funds, prime money market funds, and other counterparties. Bottom panel: the government fund share of the balance in percent. Fund holdings are repurchase agreements with the Federal Reserve Bank of New York reported on Form N-MFP, dated to the last operation of the month; the total and the non-fund remainder are the New York Fed's counterparty-type amounts for that operation. Source: SEC Form N-MFP, Federal Reserve Bank of New York, FRED series RRPONTSYD.

B The Narrative Restrictions

This appendix states the narrative restrictions of [Section 3.3](#). Each restriction ties a labeled shock to the episode that generated it through two requirements. Let W be a window of months inside a regime and j the focal shock. The first is a sign requirement on the shock's realizations. For the RFBS window the mean realization is positive,

$$\frac{1}{|W|} \sum_{t \in W} \tilde{\varepsilon}_{j,t} > 0. \quad (\text{B.1})$$

For the OFBS window the restriction reverses the inequality in (B.1). The second restricts the historical decomposition. Let $\text{HD}_{j \rightarrow L}(W)$ denote the contribution of labeled shock j to $L = \log(\text{Reserves} + \text{ON RRP})$ accumulated over the window W . The restriction requires this contribution to dominate those of the other labeled shocks,

$$|\text{HD}_{j \rightarrow L}(W)| > \sum_{\substack{i \text{ labeled} \\ i \neq j}} |\text{HD}_{i \rightarrow L}(W)|. \quad (\text{B.2})$$

In the baseline, $\text{HD}_{j \rightarrow L}(W)$ is the multi-horizon interval historical decomposition of [Antolín-Díaz and Rubio-Ramírez \(2018\)](#), the contribution of shock j to the level of $\log(\text{Reserves} + \text{ON RRP})$ at the end of the window. The contributions of the n rotated shocks add up to the unexpected movement of $\log(\text{Reserves} + \text{ON RRP})$ over the window, and each restriction assigns a share of that movement to each shock. The sum runs over the labeled shocks only, while [Antolín-Díaz and Rubio-Ramírez \(2018\)](#) compare the focal shock against every other shock.

C HARS and HARS-Z

This appendix expands [Section 4](#). Appendix [C.1](#) reviews HARS in detail. Appendix [C.2](#) derives the zero constraint matrices and states the formal consequences of adding zero restrictions to HARS. Appendix [C.3](#) builds the map γ_z . Appendix [C.4](#) states the results that establish the validity of HARS-Z. Appendix [C.5](#) states the posterior-inference estimator. The next appendix collects the proofs.

C.1 HARS in Detail

The absorbed rotation. HARS (KZ, Algorithm 1) combines likelihood-preserving posterior draws of the model parameters with rotations drawn from a well-defined probability distribution restricted by the sign and narrative restrictions. In homoskedastic SVARs, the rotations of the coefficients, $(A_0, A_+) \mapsto (A_0 Q, A_+ Q)$ with $Q \in O(n)$, are observationally equivalent. Sign restrictions identify shocks by selecting, among these rotations, those that produce economically sensible impulse responses. Under heteroskedasticity this rotation invariance fails. Rotating the coefficients while holding the regime variances fixed mixes shocks whose relative variances differ across regimes and changes the heteroskedastic likelihood (KZ, Lemma 1 and Proposition 2). The rotation would change the fitted model rather than select an interpretation.

KZ’s construction absorbs the regime scale before rotation. Define $\Xi_{s_t} := \Lambda_{s_t}^{-1/2}$ and

$$\varepsilon_t := \Xi_{s_t} \eta_t, \quad E(\varepsilon_t \varepsilon_t' \mid s_t) = I_n,$$

and the shocks have unit variance in every regime. Post-multiplying [\(1\)](#) by Ξ_{s_t} absorbs the regime scaling into the block $(A_0 \Xi_{s_t}, A_+ \Xi_{s_t})$. Rotating the entire absorbed block by $Q \in O(n)$ preserves the independent-shock representation and leaves the likelihood unchanged (KZ, Proposition 1),

$$p(y \mid \theta, Q) = p(y \mid \theta) \quad \text{for every } Q \in O(n), \quad (\text{C.1})$$

where $\theta = (A_0, \Lambda, A_+)$ collects the structural objects and $\Lambda = \{\Lambda_k\}_{k=1}^h$. Absorbing first keeps the likelihood invariant to Q under heteroskedasticity. I refer to the rotated unit-variance shocks $\tilde{\varepsilon}'_t = \varepsilon'_t Q$ as the labeled shocks. The dynamic responses of the observables to the labeled shocks in regime k at horizon ℓ are the rows of

$$\Theta_{k,\ell}(Q) = Q' \Lambda_k^{1/2} A_0^{-1} \Phi_\ell(B), \quad (\text{C.2})$$

where $\{\Phi_\ell(B)\}_{\ell \geq 0}$ denotes the moving-average coefficients of the reduced form, with $\Phi_0(B) = I_n$.²⁰ The regime affects the responses only through the positive diagonal matrix $\Lambda_k^{1/2}$. I also define the average transmission

$$\bar{\Theta}_\ell(Q) = Q' A_0^{-1} \Phi_\ell(B), \quad (\text{C.3})$$

which evaluates the responses at the average of the relative variances across regimes, and that average equals one under the normalization. The average transmission is the object I report as the impulse response. Because Q' mixes the regime-specific scales in $\Lambda_k^{1/2}$ across shocks, the regime-specific responses are not positive rescalings of the average response, and restrictions imposed regime by regime add information.

The restricted posterior target. The likelihood invariance (C.1) separates posterior simulation from identification. Write $SR(\theta, Q)$ for the indicator that the sign and narrative restrictions are satisfied, and define the admissible set $R(\theta) = \{Q \in \mathcal{O}(n) : SR(\theta, Q) > 0\}$, as in KZ. When $\mu(R(\theta)) > 0$, where μ is Haar measure on $\mathcal{O}(n)$, KZ's Definition 1 gives the restricted posterior target

$$\pi(d\theta, dQ \mid y, SR) = p(\theta \mid y) \frac{\mathbf{1}\{Q \in R(\theta)\}}{\mu(R(\theta))} \mu(dQ). \quad (\text{C.4})$$

²⁰Under the right-multiplication convention, $y'_{t+\ell} = \dots + u'_t \Phi_\ell(B)$ and $\Phi_\ell(B) = \sum_{m=1}^{\min(\ell, p)} \Phi_{\ell-m}(B) B_m$, where B_m is the m -th $n \times n$ lag block of B . Row j of $\Theta_{k,\ell}(Q)$ collects the horizon- ℓ responses of the n observables to labeled shock j . Unlike KZ, who write $\Theta_{k,\ell}(Q) = \Phi_\ell(B) Q' \Lambda_k^{1/2} A_0^{-1}$, this draft places the moving-average matrix on the right, as in (C.2), and all remaining notation is identical.

Drawing θ from its posterior and then drawing Q from Haar measure restricted to $R(\theta)$ produces exact draws from this target (KZ, Proposition 3). The restrictions truncate only the conditional distribution of the rotation. They change the support of Q , not the posterior weight on θ .

The rotation in Gaussian coordinates. HARS draws the rotation in Gaussian coordinates. An $n \times n$ matrix X with independent standard-normal entries maps to a Haar-distributed rotation $Q = \gamma(X)$ through the QR map γ . Column-wise updates of X by the elliptical slice sampler (ESS; Murray et al., 2010) leave the conditional distribution of Q in (C.4) invariant (KZ, Proposition 4). When a posterior move leaves the retained rotation stale, HARS repairs in two stages. The posterior-draw repair redraws completed parameter candidates at the fixed rotation and retains the first admissible one. This repair preserves the sign-restricted marginal target of (A_0, Λ) (KZ, Theorem 1 and Step 2a of their Algorithm 1). Only if no candidate passes does the global rotation repair draw independent rotations from Haar measure and retain the first admissible one. That draw has the restricted-Haar conditional distribution of (C.4) (KZ, Proposition 5(i) and Step 2b).

C.2 The Failure Under Zero Restrictions

Zero restrictions break this construction, and the failure originates in the measure rather than in the algorithm. The zeros invalidate the restricted target (C.4) and the Gaussian coordinatization, while the likelihood invariance (C.1) continues to hold. The obstacle is that a zero restriction is not a truncation. Sign and narrative restrictions truncate the conditional distribution of Q to a subset of positive Haar measure, and the target (C.4) remains well defined. A zero restriction instead imposes a linear equality constraint on a column of Q and lowers the dimension of the admissible set. Adding the zeros to the indicator SR fails, as Remark 1(i) states below.

Under the average transmission (C.3), the impact response of variable i_0 to shock j is $q'_j A_0^{-1} e_{i_0}$, where e_{i_0} is the i_0 -th column of I_n . The response is linear in the j -th column q_j of Q , and a zero restriction on this response is a linear constraint on q_j . For a structural draw θ , I describe the zeros on shock j with two objects, following the notation of Arias et al. (2018). The scalar z_j counts the

zero restrictions on shock j . The $z_j \times n$ matrix $Z_j(\theta)$ stacks those restrictions as rows. Each row is $(A_0^{-1} e_{i_0})'$ for an impact zero on variable i_0 and $(A_0^{-1} \Phi_\ell(B) e_{i_0})'$ for a horizon- ℓ zero, and z_j is the row count of $Z_j(\theta)$.²¹ The zero restrictions on shock j read

$$Z_j(\theta) q_j = 0,$$

and the sampler must rebuild the constraint matrices at every iteration. The matrices depend on the draw through A_0 and $\Phi_\ell(B)$.²² Define the zero manifold

$$\mathcal{Q}(\theta) = \{ Q \in \mathcal{O}(n) : Z_j(\theta) q_j = 0 \text{ for every restricted } j \}.$$

Whenever some $z_j \geq 1$, $\mathcal{Q}(\theta)$ is a lower-dimensional smooth manifold in $\mathcal{O}(n)$ and $\mu(\mathcal{Q}(\theta)) = 0$, as ARRW note in motivating their Algorithm 2.²³ The next remark states the consequences for HARS.

Remark 1 (HARS cannot impose zero restrictions). Fix θ with at least one zero restriction. Then $\mu(\mathcal{Q}(\theta)) = 0$, and the following hold.

- (i) The restricted posterior target (C.4) is not defined.
- (ii) The Gaussian coordinatization never reaches the zeros, that is, $\Pr\{\gamma(X) \in \mathcal{Q}(\theta)\} = 0$.

(iii) Every step of KZ's Algorithm 1 fails with probability one. The failure covers the posterior draw

²¹ARRW impose the zeros on a transformation $F(A_0, A_+)$ of the structural parameters, writing $Z_j F(A_0, A_+) e_j = 0$ with a constant selection matrix Z_j and e_j the j th column of I_n . In their orthogonal reduced-form parameterization the rotation carries the restriction, and their construction evaluates F before the rotation, giving $Z_j F q_j = 0$. I fold that product into a single matrix $Z_j(\theta) := Z_j F(\theta)$, which matches KZ's convention $\tilde{\varepsilon}'_t = \varepsilon'_t Q$. The fold is notational. The transformation F appears as the rows of $Z_j(\theta)$, and the count z_j , the dimension bounds, and the null-space construction are identical to theirs.

²²Unlike the admissibility check $SR(\theta, Q)$, in which the sign restrictions must hold in every regime and Λ enters, the zero constraints do not depend on Λ . The average transmission sets the relative variances to one, and the zeros hold exactly under (C.3). They need not hold in the regime-specific responses (C.2), where $\Lambda_k^{1/2}$ enters the product $Q' \Lambda_k^{1/2} A_0^{-1} \Phi_\ell(B)$ and reweights the constrained column.

²³When $Z_j(\theta)$ has full row rank, the column q_j must lie on the unit sphere of the $(n - z_j)$ -dimensional null space of $Z_j(\theta)$. The zero restrictions lower the dimension of $\mathcal{Q}(\theta)$ below the $n(n - 1)/2$ dimensions of $\mathcal{O}(n)$, and any lower-dimensional subset has Haar measure zero. ARRW establish the same fact in the structural parameterization. The set of structural parameters satisfying the sign and zero restrictions is of Lebesgue measure zero in the set of all structural parameters. The zero restrictions define a lower-dimensional smooth manifold, and the volume measure of their Theorem 3 serves as the reference measure on it.

with admissibility check (Step 1), the posterior-draw repair by the joint procedure (Step 2a), the global rotation repair (Step 2b), and the rotation draw (Step 3). The column-by-column accept–reject initialization fails as well.

Proof of Remark 1. (i) Adding the zeros to the indicator SR forces $R(\theta) \subseteq Q(\theta)$. Monotonicity gives $\mu(R(\theta)) \leq \mu(Q(\theta)) = 0$. The condition $\mu(R(\theta)) > 0$ of KZ’s Definition 1 fails at every θ with at least one zero restriction.

(ii) The map γ is the QR map that normalizes the R factor to a positive diagonal and returns the orthogonal factor. The rotation $\gamma(X)$ is Haar distributed on $O(n)$ (KZ, Proposition 4). The equality $\mu(Q(\theta)) = 0$ gives $\Pr\{\gamma(X) \in Q(\theta)\} = 0$. The construction projects each Gaussian column only onto the orthogonal complement of the preceding columns, never onto the null spaces of $Z_j(\theta)$.

(iii) The rotation-side steps of KZ’s Algorithm 1 fail by (ii). Only Steps 1 and 2a need a separate argument. Both redraw θ rather than Q .

Initialization. Each accepted column of the initialization, Algorithm 4 of [Arias et al. \(2026\)](#), is a projected Gaussian draw. A projected Gaussian draw lies in the null space of $Z_j(\theta)$ with probability zero, by (ii). Almost surely no initial admissible pair exists. The chain never occupies a state whose rotation satisfies the zeros.

Steps 1 and 2a. Both hold the retained rotation fixed and check completed draws. For a fixed unit column $q_j \neq 0$, the condition $Z_j(\theta) q_j = 0$ depends on θ only through A_0 and $\Phi_\ell(B)$. The condition is an exact alignment among functions of A_0^{-1} and $\Phi_\ell(B)$. The set of structural parameters at which the alignment holds is of Lebesgue measure zero, following the argument ARRW use for their set U (their Appendix A.3). Every completed draw and candidate comes from a proposal density and the Gibbs conditional density $p_+(A_+ \mid z_0, s_0)$. Draws with a density avoid the alignment set almost surely. When the MH step rejects and sets $z_0^\star = z_0$, the zeros already fail at the current first block, by the initialization argument.

Step 2b. Each independent Haar draw lies off $Q(\theta^\star)$ almost surely, by (ii).

Step 3. The slice step starts from a current admissible point. No such point is ever available, by the initialization argument. Even from a state on $Q(\theta)$, the update fails. Consider the ESS update of

a column j that carries zero restrictions. Following KZ's Proposition 4, draw an auxiliary vector $\tau_j \sim \mathcal{N}(0, I_n)$ and consider the ellipse

$$x_j(\vartheta) = x_j \cos \vartheta + \tau_j \sin \vartheta, \quad \vartheta \in [0, 2\pi].$$

Under γ , the proposed column is the projection of $x_j(\vartheta)$ onto the orthogonal complement of the preceding columns, normalized. Write P for that projection. The projection P is fixed along the ellipse. The preceding columns do not move when the ESS step updates column j . The normalization does not affect the zero condition. The zero restrictions hold at the proposal if and only if $Z_j(\theta) P x_j(\vartheta) = 0$. Linearity of P gives

$$Z_j(\theta) P x_j(\vartheta) = \cos \vartheta Z_j(\theta) P x_j + \sin \vartheta Z_j(\theta) P \tau_j.$$

The coefficient $Z_j(\theta) P x_j$ is zero at the current state. The current column satisfies its zeros and is proportional to $P x_j$. The coefficient $Z_j(\theta) P \tau_j$ is nonzero almost surely. The auxiliary vector is a fresh Gaussian draw, and $P \tau_j$ is a projected Gaussian column. A projected Gaussian column satisfies its zero restrictions with probability zero, by (ii). The zero restrictions hold on the ellipse only where $\sin \vartheta = 0$, at $\vartheta \in \{0, \pi\}$. At these angles the proposal is the current column up to sign. The update never reaches a new rotation on $\mathcal{Q}(\theta)$. The slice step retains only angles at which the admissibility indicator passes, and no angle off $\{0, \pi\}$ passes. $\square$

Rotations drawn from the uniform distribution satisfy the zero restrictions with probability zero.²⁴ The sampler must parameterize the zero-restricted set directly rather than reach it by sampling. It must build the rotation on $\mathcal{Q}(\theta)$, and its slice-sampling updates must never leave it.

Appendix C.3 constructs exactly such a parameterization.

²⁴The obstruction is also topological. The elliptical slice sampler of [Arias et al. \(2026\)](#) assumes an admissible set of positive measure (their Assumption 1). Strict-inequality sign restrictions give an open set of positive measure. Zero restrictions give a set that is neither open nor of positive measure.

C.3 Extending HARS to Zero Restrictions

HARS-Z extends HARS to include zero restrictions on the rotation. Developing this extension is the paper's methodological contribution. The extension replaces the QR map γ with a map $\gamma_z(\cdot; \theta)$ from a Gaussian matrix to a rotation that lies on $Q(\theta)$ *by construction*. ARRW's Algorithm 2 supplies such a map, and the zero restrictions hold exactly along the chain. The subsection builds the map γ_z and sets the state of the chain and its posterior target. It identifies the measure that the target induces on $Q(\theta)$ and constructs the rotation draw.

ARRW's Algorithm 2 builds the columns of Q one at a time, following a permutation σ of the shock indices $\{1, \dots, n\}$ that places the most-restricted shocks first. The index $\sigma(j)$ is the shock whose column the construction builds at step j . The permutation must satisfy $z_{\sigma(j)} \leq n - j$ for every j , where $z_{\sigma(j)}$ is the zero count of the shock assigned to step j . The j -th column must be orthogonal to the $j - 1$ preceding columns and must satisfy its own $z_{\sigma(j)}$ zeros. A candidate direction exists only if $(j - 1) + z_{\sigma(j)} < n$.²⁵ For every restricted shock j , the matrix $Z_j(\theta)$ must also have full row rank, that is, its z_j rows must be non-redundant.²⁶

Given a Gaussian matrix $G = [g_1, \dots, g_n]$, the construction proceeds one column at a time. At step j , the admissible directions for column $q_{\sigma(j)}$ form the linear subspace

$$N_j(\theta) := \text{null}(M_j(\theta)), \quad M_j(\theta) := \begin{bmatrix} [q_{\sigma(1)}, \dots, q_{\sigma(j-1)}]' \\ Z_{\sigma(j)}(\theta) \end{bmatrix}, \quad (\text{C.5})$$

where $M_j(\theta)$ stacks the preceding columns $q_{\sigma(1)}, \dots, q_{\sigma(j-1)}$ and the zero restrictions of shock $\sigma(j)$. The null space $N_j(\theta)$ collects the vectors orthogonal to every preceding column and consistent with those zeros.

The projection at step j requires the rows of $M_j(\theta)$ to be linearly independent. The null space has the dimension $n - (j - 1) - z_{\sigma(j)}$ that the ordering condition above presumes. Each block of $M_j(\theta)$ is independent on its own. The preceding columns are orthonormal, and $Z_{\sigma(j)}(\theta)$ has full row rank.

²⁵Later positions have more orthogonality constraints and admit fewer zeros, and the most heavily restricted shocks come first.

²⁶ARRW impose this assumption on their selection matrices Z_j when they set their notation (their Section 2.2).

The requirement binds only across the two blocks. I assume that the rows of $Z_{\sigma(j)}(\theta)$ are linearly independent of the preceding columns at every orthonormal collection $q_{\sigma(1)}, \dots, q_{\sigma(j-1)}$ whose columns satisfy their zero restrictions.²⁷ ARRW derive this property from a regularity condition on their function F and the full row rank of the Z_j . See their footnote 6 in Section 2.2 and their Appendix A.3, conditions 1 and 2. I assume it directly, at each θ the chain visits rather than over the whole parameter space.

The construction projects g_j orthogonally onto $N_j(\theta)$ and normalizes,

$$q_{\sigma(j)} = \frac{P_j(\theta) g_j}{\|P_j(\theta) g_j\|}, \quad (\text{C.6})$$

where $P_j(\theta)$ denotes the orthogonal projection onto $N_j(\theta)$ and the denominator is nonzero for almost every g_j . Write $\gamma_z(\cdot; \theta)$ for the map $G \mapsto Q$ that these n projection steps define. The chain below stores this Gaussian input as part of its state and denotes it by X .

Two properties follow. First, each column satisfies its constraints by construction, and after n steps Q is orthogonal and lies on $Q(\theta)$. The construction involves no accept–reject step. Second, the projected Gaussian column favors no direction within $N_j(\theta)$, and the normalized column is uniform over the admissible directions at step j .²⁸ This step-by-step distribution matches Steps 2 and 3 of ARRW’s Algorithm 2.²⁹

The uniform distribution of each column, conditional on the preceding columns, determines the law of the completed rotation. Write $\mu_z(\cdot; \theta)$ for the null-space construction measure on $Q(\theta)$, the

²⁷This independence fails only at draws of θ where some combination of the rows of $Z_{\sigma(j)}(\theta)$ falls exactly in the span of the preceding columns. The rows are functions of entries of A_0^{-1} and the moving-average coefficients. The failure requires an exact equality among these entries. A draw of θ from a continuous distribution satisfies an exact equality with probability zero. The retained draws avoid the failure set almost surely. ARRW make the same genericity argument for their set U (their Appendix A.3), and the implementation checks the condition at every construction step at no additional cost.

²⁸The orthogonal projection of a standard Gaussian vector onto a linear subspace is a standard Gaussian within that subspace, whose distribution is invariant under rotations of the subspace. Its normalized direction is uniform on the subspace’s unit sphere.

²⁹At each step j , they draw a standard normal vector on $\mathbb{R}^{n+1-j-z_j}$, normalize it to the unit sphere, and map it through a matrix K_j . The columns of K_j form an orthonormal basis for the null space of M_j , the matrix that stacks the preceding columns and the zero restrictions. The orthogonal projection above induces the same uniform distribution on that null space. The two constructions draw each column from the same distribution.

distribution of $\gamma_z(G; \theta)$ under a standard Gaussian matrix G ,

$$\mu_z(C; \theta) := \Pr\{\gamma_z(G; \theta) \in C\} \quad \text{for every measurable } C \subseteq \mathcal{Q}(\theta). \quad (\text{C.7})$$

Proposition 1 characterizes this measure and relates it to the conditional-Haar (uniform) measure on $\mathcal{Q}(\theta)$.

I depart from ARRW in one respect. At step j they draw $x_j \in \mathbb{R}^{n+1-j-z_j}$ from a standard normal distribution and normalize it to the unit sphere. The dimension of their coordinates varies with j and, through z_j , with the restriction pattern. I instead use a single $n \times n$ Gaussian matrix X , the coordinates in which HARS runs its ESS update. Step j takes the j -th column of X as its input. The randomness remains separate across steps, and every column keeps the fixed dimension n . The projection at each step keeps only the components in $N_j(\theta)$, and the projected column has the same distribution as the vector ARRW draw at that step. The fixed $n \times n$ matrix allows me to apply the column-wise ESS update of KZ (Proposition 4) and its validity argument without change.

Lemma 1 (Properties of the map γ_z). *Fix θ such that, for every j , the matrix $Z_{\sigma(j)}(\theta)$ has full row rank and the permutation satisfies $z_{\sigma(j)} \leq n - j$. Require also that the matrix $M_j(\theta)$ in (C.5) has full row rank at every orthonormal collection $q_{\sigma(1)}, \dots, q_{\sigma(j-1)}$ whose columns satisfy their zero restrictions. Then the following hold for almost every G .*

- (i) $\gamma_z(G; \theta) \in \mathcal{Q}(\theta)$, and every zero restriction holds exactly, with no rejection.
- (ii) Every construction step k depends on G only through $g_1, \dots, g_k$. Perturbing g_j leaves the columns from steps $1, \dots, j - 1$ unchanged, and the construction re-projects the perturbed and later columns into their own null spaces.
- (iii) If $Q \in \mathcal{Q}(\theta)$ is admissible and $G_Q := [q_{\sigma(1)}, \dots, q_{\sigma(n)}]$ collects the columns of Q in build order, then $\gamma_z(G_Q; \theta) = Q$. The construction reproduces its own output, and setting $X = G_Q$ places the chain at an admissible state.

Proof of Lemma 1. (i) At step j the subspace $N_j(\theta)$ in (C.5) is, by construction, orthogonal to

every preceding column and contained in $\text{null}(Z_{\sigma(j)}(\theta))$. The projected and normalized column $q_{\sigma(j)}$ therefore has unit norm, is orthogonal to its predecessors, and satisfies its zero restrictions exactly. The matrix $M_j(\theta)$ in (C.5) has $(j - 1) + z_{\sigma(j)}$ rows. The dimension of $N_j(\theta)$ equals $n + 1 - j - z_{\sigma(j)} \geq 1$ if and only if $M_j(\theta)$ is of full row rank (ARRW). Full row rank holds by hypothesis at every orthonormal collection of preceding columns whose columns satisfy their zero restrictions. The construction completes with Q orthogonal. The projection of a Gaussian column is zero only on a null set. The rotation $\gamma_z(G; \theta)$ lies in $Q(\theta)$ for almost every G .

(ii) Fix a step k . The subspace $N_k(\theta)$ has two ingredients, the preceding columns and the constraint matrix $Z_{\sigma(k)}(\theta)$. The constraint matrix is fixed at the given θ . The preceding columns come from steps $1, \dots, k - 1$, and step k adds g_k and nothing else. By induction, every construction step k depends on G only through $g_1, \dots, g_k$. In particular, g_j is not among the inputs of steps $1, \dots, j - 1$ and is among the inputs of steps $j, \dots, n$. Perturbing g_j leaves steps $1, \dots, j - 1$ untouched. The construction recomputes steps $j, \dots, n$. Each recomputed column passes through the projection onto its own null space. The zeros hold after the perturbation.

(iii) Let $Q \in Q(\theta)$ be admissible. Each of its columns already lies in its own subspace $N_j(\theta)$. Orthogonal projection acts as the identity on $N_j(\theta)$. Normalization preserves a unit vector. The construction reproduces Q column by column. The reproduction property is the projection argument ARRW use to establish that their Algorithm 2 can generate every rotation satisfying the zero restrictions. $\square$

Lemma 1 makes ARRW's Algorithm 2 usable inside a Markov chain. Part (i) restates the property of their Algorithm 2 in the fixed $n \times n$ Gaussian parameterization introduced above. Part (ii) is new. ARRW draw each rotation independently and never perturb one afterward. They never need to ask how the rotation changes when a single Gaussian column moves.³⁰ The column-wise ESS update moves one column of X along an ellipse and poses exactly that question. Part (ii) supplies the answer. The recomputed rotation stays on $Q(\theta)$. Part (iii) reuses ARRW's support argument

³⁰ARRW use their coordinates only for independent proposals. The QR map γ of KZ fails for the reason given in Remark 1(ii). Under γ , an elliptical perturbation of a single Gaussian column cannot restore the zeros, and the sign-only ESS never visits the manifold.

for a new purpose. Setting $X = G_Q$ starts the chain at any admissible rotation, whatever search produced that rotation. Part (iii) supplies the admissible state that [Algorithm 1](#) presupposes.

The state of the chain. A sampler could store the rotation itself, keeping the matrix Q in the state and reusing it after every structural update. HARS does exactly this. KZ write the state as (θ, Q) and set $X = Q$ whenever the ESS step needs Gaussian coordinates. The choice is harmless there. The QR map γ does not depend on θ , and the stored rotation and its coordinates never disagree. On $Q(\theta)$ the choice matters. The manifold $Q(\theta)$ moves with θ . A stored Q lies on the old manifold and leaves the new manifold $Q(\theta^*)$ with probability one as soon as the structural draw moves. The zero restrictions fail after every structural update. I store the Gaussian coordinates instead. The state is the pair (θ, X) , where X is an $n \times n$ Gaussian matrix, and the sampler recomputes the rotation from the state whenever it needs one,

$$Q = \gamma_z(X; \theta).$$

The next subsection establishes the validity of the construction. [Proposition 1](#) derives the measure that the construction induces on the admissible set, [Corollary 2](#) shows that the weighted draws recover conditional-Haar expectations, and [Proposition 2](#) gives the column-wise elliptical slice sampling update that the sampler uses. [Algorithm 1](#) collects the steps.

C.4 The Validity of HARS-Z

A structural update leaves X unchanged and changes only the rotation $\gamma_z(X; \theta)$. The construction projects the same Gaussian columns onto the null spaces of the new draw, and $\gamma_z(X; \theta^*)$ lies on $Q(\theta^*)$ by [Lemma 1\(i\)](#). The zeros hold after every structural move by construction, and the admissibility check reduces to the sign and narrative restrictions. The rotation remains auxiliary, exactly as in KZ. Given (θ, X) , the rotated absorbed system of [\(C.1\)](#) evaluated at $Q = \gamma_z(X; \theta)$ generates the data, and the conditional density $p(y \mid \theta, X)$ is well defined. The next definition states the posterior target of HARS-Z, and the corollary that follows establishes its validity.

Definition 1 (The restricted posterior target on (θ, X)). Let $\varphi(X)$ denote the standard Gaussian density of the $n \times n$ matrix X and define

$$\kappa(\theta) = \int \mathbf{1}\{SR(\theta, \gamma_z(X; \theta)) > 0\} \varphi(X) dX,$$

the Gaussian probability of the admissible set at θ . When $\kappa(\theta) > 0$, define the conditional distribution of X given θ ,

$$\pi_X(dX \mid \theta, SR) := \frac{\mathbf{1}\{SR(\theta, \gamma_z(X; \theta)) > 0\} \varphi(X)}{\kappa(\theta)} dX. \quad (\text{C.8})$$

The restricted posterior target is

$$\pi(d\theta, dX \mid y, SR) = \pi(\theta \mid y) \pi_X(dX \mid \theta, SR) d\theta. \quad (\text{C.9})$$

[Definition 1](#) is KZ's Definition 1 in Gaussian coordinates, with $\kappa(\theta) > 0$ replacing $\mu(R(\theta)) > 0$. KZ's definition takes one fixed distribution, Haar measure on $O(n)$, and restricts it to the θ -dependent region $R(\theta)$. That structure is unavailable on rotations here. A distribution on $Q(\theta)$ assigns no mass to $Q(\theta')$ at any other draw θ' , and no fixed distribution on rotations exists. The Gaussian distribution of X restores the structure. Equation (C.8) restricts the same $\varphi(X)$ to the admissible set at every draw, and θ enters only through the indicator and the map $\gamma_z(\cdot; \theta)$.

Corollary 1 (The Gaussian coordinates are auxiliary). *Fix θ satisfying the conditions of [Lemma 1](#) and write $Q = \gamma_z(X; \theta)$. For almost every X ,*

$$p(y \mid \theta, X) = p(y \mid \theta, Q) = p(y \mid \theta).$$

If, in addition, $\kappa(\theta) > 0$, then drawing θ from its posterior and then drawing X from the conditional $\pi_X(\cdot \mid \theta, SR)$ in (C.8) produces exact draws from the joint target (C.9). Sampling X in a separate block leaves the marginal posterior of θ unchanged.

Proof of [Corollary 1](#). Fix θ and the regime path. Given (θ, X) , the rotated absorbed system

evaluated at $Q = \gamma_z(X; \theta)$ generates the data,

$$y_t' A_0 \Xi_{s_t} Q = x_t' A_+ \Xi_{s_t} Q + \tilde{\varepsilon}_t', \quad E(\tilde{\varepsilon}_t \tilde{\varepsilon}_t' | s_t) = I_n,$$

where $\tilde{\varepsilon}_t' = \varepsilon_t' Q = \eta_t' \Xi_{s_t} Q$ collects the labeled shocks of Appendix C.1. The conditional density of y given (θ, X) is $p(y | \theta, Q)$ evaluated at $Q = \gamma_z(X; \theta)$. The evaluation gives the first equality of the corollary. By Lemma 1(i), $\gamma_z(X; \theta) \in \mathcal{Q}(\theta) \subseteq \mathcal{O}(n)$ for almost every X . KZ's Proposition 1 gives $p(y | \theta, Q) = p(y | \theta)$ for every $Q \in \mathcal{O}(n)$. Evaluating at $Q = \gamma_z(X; \theta)$ yields the second equality of the corollary. The dependence of the map on θ does not affect the evaluation. The invariance $p(y | \theta, Q) = p(y | \theta)$ holds at each fixed θ for every $Q \in \mathcal{O}(n)$, and $\gamma_z(X; \theta)$ is one such rotation.

For the last claim, let $\kappa(\theta) > 0$. The conditional $\pi_X(\cdot | \theta, SR)$ of (C.8) is then well defined. The conditional is the Gaussian distribution of X restricted to the admissible set $\{X : SR(\theta, \gamma_z(X; \theta)) > 0\}$. Specify the prior of the pair as $p(\theta)$ for θ and $\pi_X(\cdot | \theta, SR)$ for X given θ . The joint posterior under this prior satisfies

$$\begin{aligned} \pi(\theta, X | y) &\propto p(y | \theta, X) p(\theta) \pi_X(X | \theta, SR) \\ &= p(y | \theta) p(\theta) \pi_X(X | \theta, SR) \\ &\propto \pi(\theta | y) \pi_X(X | \theta, SR). \end{aligned}$$

The first line is Bayes' rule. The second line applies the first part. The third line collects $\pi(\theta | y) \propto p(y | \theta) p(\theta)$. The conditional distribution of X is the Gaussian distribution restricted to the admissible set. The posterior density of θ remains $\pi(\theta | y)$. The joint posterior under this prior is exactly the target (C.9). Drawing θ from its posterior and then drawing X from $\pi_X(\cdot | \theta, SR)$ produces exact draws from the target. The argument is KZ's Proposition 3 with Haar measure restricted to $R(\theta)$ replaced by the Gaussian distribution of X restricted to the admissible set. $\square$

Corollary 1 justifies the two-block structure of the sampler. The equalities say that the likelihood never involves X . The θ block runs the posterior sampler of BPSS exactly as in KZ. The last claim

covers the rotation block. Drawing the rotation at fixed θ from the conditional (C.8) produces draws from the target (C.9) and leaves the marginal posterior of θ unchanged.

Posterior inference evaluates the impulse responses, functions of (θ, Q) . The next definition states the distribution that posterior inference targets.

Definition 2 (The conditional-Haar target on (θ, Q)). Fix the conditions of Lemma 1 at every θ with $\kappa(\theta) > 0$. The conditional-Haar target, restated from Section 4, is

$$\pi_v(d\theta, dQ \mid y, SR) = \pi(\theta \mid y) \frac{\mathbf{1}\{Q \in R(\theta)\} \mu(dQ \mid Q(\theta))}{\mu(R(\theta) \mid Q(\theta))} d\theta, \quad (6)$$

where $\mu(\cdot \mid Q(\theta))$ denotes the conditional-Haar (uniform) measure on $Q(\theta)$ and $R(\theta) = \{Q \in Q(\theta) : SR(\theta, Q) > 0\}$ denotes the admissible set on the zero-restricted set.

Remark 2. The normalizer of (6) is positive whenever $\kappa(\theta) > 0$. The admissible set $R(\theta)$ is relatively open in $Q(\theta)$, as the restrictions hold with strict inequalities. When $R(\theta)$ is empty, no draw satisfies the indicator in (C.8) and $\kappa(\theta) = 0$. The uniform measure on $Q(\theta)$ assigns positive mass to every nonempty relatively open subset.

Definition 2 is the analogue of KZ's target (C.4), with Haar measure on $O(n)$ replaced by the conditional-Haar measure on $Q(\theta)$. The restrictions truncate only the conditional distribution of the rotation, as in Proposition 3 of KZ, and the marginal posterior of θ remains $\pi(\theta \mid y)$, consistent with Corollary 1. Definition 1 states the target on (θ, X) , the coordinates in which the chain runs, and Definition 2 states the target on (θ, Q) , the coordinates of the impulse responses.

Proposition 1 (The induced measure on $Q(\theta)$). Fix θ satisfying the conditions of Lemma 1 with $\kappa(\theta) > 0$. The distribution over $Q(\theta)$ that $\pi_X(\cdot \mid \theta, SR)$ induces through γ_z is the null-space construction measure $\mu_z(\cdot; \theta)$ of (C.7), restricted to the admissible set $R(\theta)$. The conditional of (6) given θ is the conditional-Haar measure restricted to $R(\theta)$, written $\pi_v(dQ \mid \theta, y, SR)$. The restricted construction measure has two properties.

- (a) When the zeros restrict a single column of Q and the ordering builds that column first, the restricted construction measure coincides with $\pi_v(\cdot \mid \theta, y, SR)$.

(b) Otherwise the restricted construction measure differs from $\pi_v(\cdot \mid \theta, y, SR)$ by the volume element $v(Q; \theta)$ of the construction (ARRW, Theorem 3), and weighting each retained draw by $v(Q; \theta)$ recovers $\pi_v(\cdot \mid \theta, y, SR)$.

Proof of Proposition 1. Two preliminary steps characterize the induced distribution. Parts (a) and (b) then follow in turn.

Step 1: π_X induces the construction measure restricted to the admissible set. The argument follows the image claim of KZ's Proposition 4, with the QR map γ replaced by γ_z and Haar measure replaced by the construction measure $\mu_z(\cdot; \theta)$. Recall the conditional (C.8),

$$\pi_X(dX \mid \theta, SR) = \frac{\mathbf{1}\{SR(\theta, \gamma_z(X; \theta)) > 0\} \varphi(X)}{\kappa(\theta)} dX.$$

Without the indicator the conditional reduces to the Gaussian density φ . The matrix X then has the same distribution as the Gaussian matrix G in (C.7). The relation $\gamma_z(X; \theta) \sim \mu_z(\cdot; \theta)$ holds by definition.

Now restore the indicator. The indicator in (C.8) does not test X directly. The indicator builds the rotation $\gamma_z(X; \theta)$ and tests whether that rotation satisfies the sign and narrative restrictions. The input X passes exactly when its rotation lies in the admissible set $R(\theta) = \{Q \in O(n) : SR(\theta, Q) > 0\}$ of Section 4,

$$\{X : SR(\theta, \gamma_z(X; \theta)) > 0\} = \{X : \gamma_z(X; \theta) \in R(\theta)\}. \quad (\text{C.10})$$

Take any set C of rotations. Under the conditional $\pi_X(\cdot \mid \theta, SR)$, the probability that the rotation lands in C is

$$\begin{aligned} \Pr\{\gamma_z(X; \theta) \in C \mid SR(\theta, \gamma_z(X; \theta)) > 0\} &= \Pr\{\gamma_z(X; \theta) \in C \mid \gamma_z(X; \theta) \in R(\theta)\} \\ &= \frac{\Pr\{\gamma_z(X; \theta) \in C \cap R(\theta)\}}{\Pr\{\gamma_z(X; \theta) \in R(\theta)\}} \\ &= \frac{\mu_z(C \cap R(\theta); \theta)}{\mu_z(R(\theta); \theta)}. \end{aligned}$$

The first line rewrites the conditioning event. Equation (C.10) shows the two events are the same set of X . The second line applies the definition of conditional probability. The probabilities on the second line are unconditional. On the second line $X \sim \varphi$, and the relation $\gamma_z(X; \theta) \sim \mu_z(\cdot; \theta)$ from the start of the proof evaluates each probability as a μ_z -value on the third line.

The set C is arbitrary, and the values on every set determine a distribution. The induced distribution is

$$\frac{\mathbf{1}\{SR(\theta, Q) > 0\} \mu_z(dQ; \theta)}{\mu_z(R(\theta); \theta)}. \quad (\text{C.11})$$

The measure (C.11) is the construction measure restricted to the admissible set $R(\theta)$. The restricted construction measure is the counterpart of KZ's Haar measure restricted to $R(\theta)$.

The normalizer of (C.11) coincides with the normalizer $\kappa(\theta)$ of (C.8), $\kappa(\theta) = \mu_z(R(\theta); \theta)$. Both equal the probability of one event, that the rotation is admissible. The normalizer $\kappa(\theta)$ computes this probability in X . The normalizer $\mu_z(R(\theta); \theta)$ computes it on $Q(\theta)$.

Step 2: relation to conditional-Haar. The comparison proceeds column by column. The construction builds the rotation column by column. The construction generates $q_{\sigma(1)}, \dots, q_{\sigma(n)}$ in the order σ . Each step projects a fresh Gaussian column onto the null space $N_j(\theta)$ that the zero restrictions and the previously built columns determine, then normalizes. The projected Gaussian column favors no direction within $N_j(\theta)$. Conditional on the preceding columns,

$$q_{\sigma(j)} \mid q_{\sigma(1)}, \dots, q_{\sigma(j-1)} \sim \text{uniform on the unit sphere of } N_j(\theta).$$

The uniform law of the projected column is the second property of the construction in Appendix C.3. Write $\mu_j(\cdot \mid q_{\sigma(1)}, \dots, q_{\sigma(j-1)}; \theta)$ for this conditional distribution, the uniform distribution on the unit sphere of $N_j(\theta)$. The chain rule of probability writes the joint law as the product of these conditionals,

$$\mu_z(dQ; \theta) = \prod_{j=1}^n \mu_j(dq_{\sigma(j)} \mid q_{\sigma(1)}, \dots, q_{\sigma(j-1)}; \theta). \quad (\text{C.12})$$

Equation (C.7) defines $\mu_z(\cdot; \theta)$ only as the distribution of $\gamma_z(G; \theta)$. The product (C.12) makes the

measure explicit. Parts (a) and (b) below use (C.12).

Part (a): the zeros load on a single column. The ordering builds the restricted column first. The first null space $N_1(\theta)$ then equals the null space of the zero restrictions. Every later null space $N_j(\theta)$, $j \geq 2$, is the orthogonal complement of the preceding columns and carries no further constraint.

The comparison involves two distinct measures on $\mathcal{Q}(\theta)$. The construction measure $\mu_z(\cdot; \theta)$ is the law the sampler induces. The measure $\mu(\cdot \mid \mathcal{Q}(\theta))$ is the conditional-Haar (uniform) measure on $\mathcal{Q}(\theta)$ and serves as the reference measure of the target. The subscript z denotes the measure that the zero-restricted construction induces. The conditioning on $\mathcal{Q}(\theta)$ denotes Haar measure restricted to $\mathcal{Q}(\theta)$. Factor both measures at the first column,

$$\mu_z(dQ; \theta) = \mu_1(dq_{\sigma(1)}; \theta) \mu_z(dq_{\sigma(2)}, \dots, dq_{\sigma(n)} \mid q_{\sigma(1)}; \theta), \quad (\text{C.13})$$

$$\mu(dQ \mid \mathcal{Q}(\theta)) = \mu(dq_{\sigma(1)} \mid \mathcal{Q}(\theta)) \mu(dq_{\sigma(2)}, \dots, dq_{\sigma(n)} \mid q_{\sigma(1)}, \mathcal{Q}(\theta)), \quad (\text{C.14})$$

and compare factor by factor.

The first factors of (C.13) and (C.14) coincide. The first factor of (C.13) is $\mu_1(\cdot; \theta)$, the uniform distribution on the unit sphere of $N_1(\theta)$, by the second property of the construction. The first factor of (C.14) is uniform on the same sphere, by the following symmetry argument. Take any orthogonal matrix R that rotates $N_1(\theta)$ within itself and fixes its orthogonal complement. The map $Q \mapsto RQ$ moves the first column along the sphere and leaves every zero restriction intact. The map $Q \mapsto RQ$ therefore sends $\mathcal{Q}(\theta)$ onto itself. Haar measure is invariant under this map. Conditional-Haar inherits the invariance. Every rotation of $N_1(\theta)$ leaves the marginal distribution of $q_{\sigma(1)}$ unchanged. The only probability distribution on a sphere that every rotation leaves unchanged is the uniform distribution. Both first factors are the same uniform law on the same sphere.

The second factors of (C.13) and (C.14) coincide as well. Fix $q_{\sigma(1)}$. No zero restriction remains after the first column. Every later null space is the orthogonal complement of the preceding columns inside $q_{\sigma(1)}^\perp$. Each remaining step projects a fresh Gaussian column onto that complement and normalizes. A projected standard Gaussian column is standard Gaussian within the complement.

The remaining steps therefore run the QR map γ of KZ on $q_{\sigma(1)}^\perp$. The rotation $\gamma(X)$ of a Gaussian matrix is Haar distributed, by the Gaussian-to-Haar construction of [Rubio-Ramírez et al. \(2010\)](#) in KZ’s Proposition 4. The second factor of (C.13) is the uniform distribution over the orthonormal completions of $q_{\sigma(1)}$.

The second factor of (C.14) is the same uniform distribution. Apply the symmetry argument of the first factor to rotations that fix $q_{\sigma(1)}$ and rotate its complement. Every such rotation preserves $Q(\theta)$ and preserves conditional-Haar. The conditional law of the completions is rotation invariant. The uniform distribution is the only rotation-invariant law. ARRW report constant importance weights for their Algorithm 3 without zero restrictions, consistent with the equality of the two factors.

The factorizations (C.13) and (C.14) coincide factor by factor. The two measures are equal,

$$\mu_z(\cdot; \theta) = \mu(\cdot \mid Q(\theta)).$$

The restrictions of equal measures to $R(\theta)$ are equal. The volume element is constant. This proves part (a).

Part (b): the zeros load on two or more columns. The construction still draws each column uniformly on the unit sphere of its null space, as in (C.12). The measure $\mu_z(\cdot; \theta)$ is the law of the construction that ARRW’s Theorem 3 describes. Their change-of-variable theorem for maps into lower-dimensional smooth manifolds gives $\mu_z(\cdot; \theta)$ a density proportional to $1/v(Q; \theta)$ with respect to the uniform (volume) measure on $Q(\theta)$. The factor $v(Q; \theta)$ is the volume element of the map implied by their Algorithm 2 at fixed θ . The ratio of the two densities gives

$$\frac{d\mu(\cdot \mid Q(\theta))}{d\mu_z(\cdot; \theta)}(Q) \propto v(Q; \theta).$$

The density relation holds on all of $Q(\theta)$ and hence on $R(\theta)$. Weighting each retained draw by $v(Q; \theta)$ recovers the conditional-Haar measure restricted to $R(\theta)$, the conditional $\pi_v(\cdot \mid \theta, y, SR)$ of (6). The weighting is necessary. As ARRW document, the distribution that their Algorithm 2

implies is not invariant to a reordering of the shocks. This proves part (b) and the proposition. $\square$

Proposition 1 states that the conditional $\pi_X(\cdot \mid \theta, SR)$ induces the construction measure on $Q(\theta)$ through γ_z , not the uniform (conditional-Haar) measure on $Q(\theta)$. When the zeros load on a single column and the ordering builds that column first, the two coincide and need no correction. With zeros on more than one column, I weight the retained rotations by the volume element and normalize the weights within each retained draw of θ .³¹ **Proposition 1**(b) states the correction as a density, known only up to a constant that depends on θ . The next corollary removes the constant. Ratios of weighted expectations under the sampled measure reproduce conditional expectations under (6) exactly.

Corollary 2 (Conditional-Haar expectations from weighted draws). *Fix θ satisfying the conditions of Lemma 1 with $\kappa(\theta) > 0$. Let $E_z[\cdot \mid \theta]$ denote the expectation under the restricted construction measure of Proposition 1, the distribution of the retained rotations at θ . For every function h integrable under $\pi_v(\cdot \mid \theta, y, SR)$,*

$$\int h(\theta, Q) \pi_v(dQ \mid \theta, y, SR) = \frac{E_z[v(Q; \theta) h(\theta, Q) \mid \theta]}{E_z[v(Q; \theta) \mid \theta]}. \quad (\text{C.15})$$

Proof of Corollary 2. Write $\tilde{\mu}_z(dQ \mid \theta)$ for the restricted construction measure (C.11). By Proposition 1(b), the conditional of (6) has a density proportional to $v(Q; \theta)$,

$$\pi_v(dQ \mid \theta, y, SR) = c(\theta) v(Q; \theta) \tilde{\mu}_z(dQ \mid \theta),$$

³¹The weight is the part of the importance weight of ARRW's Algorithm 3 that survives at fixed θ . Their importance weight is the ratio of the target normal-generalized-normal density to the density of their Algorithm 2, a product of two volume elements. The element v_{f_h} converts between the structural and the orthogonal reduced-form parameterizations (their Proposition 1). The element $v_{(g \circ f_h) \mid Z}$ corrects their Algorithm 2 on the zero-restricted set (their Theorem 3). Only the second element arises here. The Metropolis step targets the structural posterior of θ directly, and the sampled and target distributions differ only in the conditional of Q given θ . The volume element $v(Q; \theta)$ of Proposition 1 is their $v_{(g \circ f_h) \mid Z}$ at fixed θ .

with a constant $c(\theta) > 0$. For every h integrable under the conditional,

$$\int h(\theta, Q) \pi_v(dQ \mid \theta, y, SR) = c(\theta) \int h(\theta, Q) v(Q; \theta) \tilde{\mu}_z(dQ \mid \theta) = c(\theta) E_z[v(Q; \theta) h(\theta, Q) \mid \theta].$$

The choice $h = 1$ makes the left side one and pins down the constant, $c(\theta) = 1/E_z[v(Q; \theta) \mid \theta]$.

Substituting the constant into the display gives (C.15). $\square$

The rotation draw. With the target of Definition 1 in place, the rotation draw keeps θ fixed and applies the column-wise ESS update of KZ (Proposition 4) through γ_z . The next proposition establishes its validity on $Q(\theta)$.

Proposition 2 (Column-wise ESS on $Q(\theta)$). *Fix θ satisfying the conditions of Lemma 1 with $\kappa(\theta) > 0$, and let $\pi_X(\cdot \mid \theta, SR)$ denote the conditional distribution of X given θ in (C.8). The column-wise ESS update proposes*

$$x_j(\vartheta) = x_j \cos \vartheta + \tau_j \sin \vartheta, \quad \tau_j \sim \mathcal{N}(0, I_n),$$

where x_j denotes the j -th column of the current $X = [x_1, \dots, x_n]$, and checks admissibility through γ_z . It has the following properties.

- (a) *Almost surely, it produces proposals whose rotations satisfy every zero restriction exactly along the entire ellipse, and the zeros never reject a proposal.*
- (b) *Each column update terminates almost surely.*
- (c) *It leaves $\pi_X(\cdot \mid \theta, SR)$ invariant.*

Proof of Proposition 2. We prove the three parts in turn.

Part (a): the zeros never reject a proposal. Fix a column j , and let $X^*(\vartheta)$ denote the current X with column j replaced by $x_j(\vartheta)$. The sampler evaluates admissibility only through the candidate rotation $\gamma_z(X^*(\vartheta); \theta)$. By Lemma 1(i), every candidate satisfies the zero restrictions exactly wherever the

construction is defined. The construction fails only where some projected column in (C.6) equals the zero vector, an exceptional set of inputs with Gaussian probability zero.

Lemma 1(i) concerns a fresh Gaussian input. The candidates along the ellipse are not fresh Gaussian inputs. A Fubini argument extends the statement to the ellipse. Fix an angle with $\sin \vartheta \neq 0$. Given the current state, the candidate column $x_j(\vartheta) = x_j \cos \vartheta + \tau_j \sin \vartheta$ is a full-support Gaussian draw, and the remaining columns arise from draws with densities. The candidate matrix $X^\star(\vartheta)$ lies in the exceptional set with probability zero. By Fubini's theorem over the uniformly drawn angle, with nonnegative integrands, the exceptional angles form a null set almost surely. The zeros hold along the entire ellipse off this null set, and the zeros never reject a proposal. The slice step retains only angles at which the sign and narrative indicator holds, exactly as in KZ's HARS.

Part (b): almost-sure termination. We argue in three steps.

Step 1: the angle $\vartheta = 0$ is admissible. The bracket of the slice search always contains $\vartheta = 0$. At $\vartheta = 0$ the candidate matrix is the current X . By **Lemma 1**(iii), γ_z maps the current X to the current rotation. The chain retains only admissible states, and the current rotation is admissible.

Step 2: $\vartheta \mapsto \gamma_z(X^\star(\vartheta); \theta)$ is continuous at $\vartheta = 0$ almost surely. Each construction step depends on the stacked matrix $M_j(\theta)$ in (C.5), abbreviated M_j , only through the orthogonal projector P_j onto its null space $N_j(\theta)$. The projector P_j does not depend on the choice of orthonormal basis. Wherever M_j has full row rank,

$$P_j = I_n - M_j'(M_j M_j')^{-1} M_j, \quad (\text{C.16})$$

a continuous function of M_j . The proposition fixes θ under the maintained conditions of **Lemma 1**. Those conditions give M_j full row rank at every orthonormal collection of preceding columns whose columns satisfy their zero restrictions. Along the ellipse, the construction steps $1, \dots, j-1$ do not move, and only the construction steps from j onward vary, by **Lemma 1**(ii). Off the exceptional set of part (a), each construction step is the composition of $M_j \mapsto P_j$ through (C.16) and the normalization in (C.6). Both pieces are continuous at almost every angle. The composition of the n

steps is continuous at $\vartheta = 0$ almost surely.

Step 3: the bracket-shrinking rule terminates. The sign and narrative restrictions are strict inequalities. The admissible set of rotations is open. By Steps 1 and 2, the admissible angular set contains an open neighborhood of $\vartheta = 0$ and is nonempty. The bracket-shrinking rule reaches the admissible angular set in finitely many steps almost surely. Steps 1 and 3 make explicit the nonemptiness assertion in the proof of KZ’s Proposition 4. That assertion suffices in their setting. The QR map γ is continuous at every full-rank X . Step 2 supplies the analogous continuity for γ_z , the only new ingredient.

Part (c): invariance. The target (C.8) is a standard Gaussian density multiplied by a nonnegative likelihood factor, here the admissibility indicator. The Gaussian density factorizes across columns, $\varphi(X) \propto \exp(-\frac{1}{2} \sum_{j=1}^n \|x_j\|^2)$, and the full conditional of a single column is

$$p(x_j \mid x_{-j}, \theta, SR) \propto \mathbf{1}\{SR(\theta, \gamma_z(X; \theta)) > 0\} \exp(-\frac{1}{2} \|x_j\|^2),$$

a Gaussian density times an indicator, as in KZ’s Proposition 4. Elliptical slice sampling is exact for targets of this form (Murray et al., 2010). Part (a) reduces the indicator along every ellipse to a sign-and-narrative indicator. Part (b) supplies the almost-sure termination. The argument of KZ’s Proposition 4 applies. Each column update samples from the admissible angular set and satisfies detailed balance for the full conditional $p(x_j \mid x_{-j}, \theta, SR)$. The systematic scan over the columns composes invariant full-conditional updates. The scan leaves $\pi_X(\cdot \mid \theta, SR)$ invariant. The retained states are serially dependent draws from $\pi_X(\cdot \mid \theta, SR)$. This proves the proposition. $\square$

Lemma 1 (ii) permits the column-wise ESS update on $Q(\theta)$. The zeros hold under every elliptical proposal by construction, and the slice check evaluates only sign and narrative admissibility. The angle-zero proposal is the current admissible state. The admissible angular set is nonempty, and the bracket-shrinking rule terminates (Proposition 2(b)).

The validity of Algorithm 1 step by step. The validity of HARS-Z follows step by step from KZ. The posterior draw in Step 1 is the sampler of BPSS, and its invariant distribution is the posterior target of θ . By Proposition 2, the rotation draw in Step 3 leaves the X -block conditional target (C.8) invariant. The retained rotations of Step 3 are serially dependent but are valid draws from the Q -block conditional target (KZ, Definition 1 and Proposition 4), extended to $Q(\theta)$ by Proposition 2. Their distribution at fixed θ is the induced measure of Proposition 1. The global rotation repair in Step 2b targets the same restricted construction measure through independent proposals. By KZ's Theorem 1, the posterior-draw repair in Step 2a redraws a completed parameter candidate and preserves the sign-restricted marginal target of $z_0 = (A_0, \Lambda)$. The sampler never uses the conditional shortcut for z_+ , which would have led to the incorrect marginal target of z_0 .³² Because X does not enter the likelihood, by Corollary 1, none of the rotation rebuilds, rotation draws, or repair steps reweights θ . The retained pairs (θ, X) are draws from the target of Definition 1. Where the zeros load on more than one column, the volume weight of Proposition 1(b) restores the conditional $\pi_v(\cdot \mid \theta, y, SR)$.

C.5 Posterior Inference

Let θ_d denote the d -th retained draw of θ , $Q^{(d,r)} = \gamma_z(X^{(d,r)}; \theta_d)$ its r -th retained rotation, and $v_{d,r} = v(Q^{(d,r)}; \theta_d)$ the volume element at that rotation. The estimator of the posterior mean of a function h of the impulse responses is

$$\widehat{E}[h] = \frac{1}{D} \sum_{d=1}^D \frac{\sum_{r=1}^{N_{\text{ess}}} v_{d,r} h(\theta_d, Q^{(d,r)})}{\sum_{r=1}^{N_{\text{ess}}} v_{d,r}}, \quad (\text{C.17})$$

where D is the number of retained draws of θ . The reported quantiles apply the same weights, normalized within each retained draw. The numerator and the denominator of the inner ratio are

³²The natural alternative holds $z_0 = (A_0, \Lambda)$ from the stale draw fixed and redraws only $z_+ = A_+$ until the check passes. This targets the wrong distribution. Under the sign-restricted posterior, the marginal of z_0 has the admissibility weight $\omega_0(z_0; s_0, Q)$, the probability that a fresh Gibbs draw of A_+ completes the pair admissibly. Redrawing A_+ until success gives every z_0 a completion. It strips this factor and over-represents structural values whose admissible completions are rare (KZ, Theorem 1).

sample averages of the numerator and the denominator on the right side of (C.15) at θ_d . The outer average runs over the posterior draws of θ , whose marginal the rotation block leaves unchanged by Corollary 1. The estimator is consistent for expectations under (6) as D and N_{ess} grow. Any θ -dependent rescaling of the volume element cancels in the inner ratio. In the single-column case of Proposition 1(a), the weights are constant within each draw and (C.17) reduces to the unweighted average of KZ.

D Estimation Diagnostics

The baseline estimates come from one chain of 10,000 retained draws after 1,000 burn-in iterations, with $N_{\text{ess}} = 5$ rotations kept per sweep. The stability mask keeps 9,615 of the draws. The repair budgets are $N_{\theta} = 300$ and $N_Q = 200$.

Sampler diagnostics. The Metropolis acceptance rate is 48 percent. The chain discards the remaining 52 percent of the proposed transitions, the stale rotations that the joint redraw does not repair within the budgets. The joint structural repair succeeds on 1,050 of 372,322 attempts.

The volume weights stay near uniform, with a within-block weight ESS share of 0.997.

Residual diagnostics. Figure D.1 plots the five equations Figure 4 leaves out. Table D.1 repeats that estimation with the four variance regimes of Table 1 in place of the single regime and standardizes each structural residual by the square root of its estimated regime variance λ_{i,s_t} . Under Gaussian errors these standardized residuals would show excess kurtosis near zero.³³ Pooled over the eight equations and the full sample they show 5.0, and the SOMA equation alone shows 24.1.

The two-step procedure of Brunnermeier et al. (2021) fits 5.84 degrees of freedom to these residuals, and the estimation uses that value. The fitted distribution matches the residuals on how often extreme residuals occur. In a sample of 1,648 observations it puts the expected number of residuals beyond six standard deviations at 0.6, and the sample holds one. Beyond five standard deviations it expects 1.5, and the sample holds two.

Table D.2 reports the estimated variance ratios $\omega_{i,s} = \lambda_{i,s}/\lambda_{i,1}$ over the retained draws.

The regime profiles differ across equations. The one-year yield variance stays above its Regime 1 level in all three later regimes. The SOMA variance rises elevenfold in Regime 3 and falls to a tenth of Regime 1 in the other two. Those two equations also carry the widest bands, spanning a factor of two or more in at least one regime. The bands of the other six stay inside a factor of 1.7 everywhere.

³³Kurtosis is the expectation of the fourth power of a standardized random variable and equals three for a Gaussian. Excess kurtosis subtracts this benchmark. Positive values indicate that extreme observations occur more often than under a Gaussian. A Student- t with $\delta > 4$ degrees of freedom has excess kurtosis $6/(\delta - 4)$.




Figure D.1: Reduced-form residuals under a single variance regime, the remaining five equations.
Note. The equations [Figure 4](#) leaves out, on the same estimation and with the same bands and subsample standard deviations. Units are percent for the log levels and percentage points for the ratio and the EBP.

Table D.1: Regime-standardized residuals under four variance regimes.

	Excess kurtosis	$\max z $
One-year yield	3.7	5.5
SOMA	24.1	9.6
log(Reserves + ON RRP)	0.6	3.5
ON RRP ratio	1.1	3.7
GDP	1.5	3.8
GDPDEF	1.3	3.8
Term spread	0.6	3.5
EBP	1.5	4.2
Pooled	5.0	9.6

Note: The table estimates equation (1) with the four variance regimes of [Table 1](#), Gaussian errors, no identification restrictions, and the prior of [Section 3.1](#). It reports the excess kurtosis and the largest absolute value of the structural residuals standardized by $\sqrt{\lambda_{i,s_t}}$. Entries are posterior medians. The pooled row stacks all equations over the full sample.

Table D.2: Estimated variance ratios by regime, $\omega_{i,s} = \lambda_{i,s}/\lambda_{i,1}$.

Equation	Regime 2	Regime 3	Regime 4
One-year yield	11.95 [7.93, 16.26]	65.75 [45.29, 89.55]	49.38 [34.37, 68.13]
SOMA	0.10 [0.06, 0.14]	11.22 [10.60, 12.50]	0.11 [0.07, 0.18]
log(R + RRP)	0.76 [0.75, 0.77]	0.95 [0.94, 0.95]	0.89 [0.86, 0.92]
ON RRP ratio	3.77 [3.46, 4.43]	5.75 [5.23, 6.77]	4.14 [3.76, 4.83]
GDP	0.37 [0.35, 0.40]	1.78 [1.73, 1.82]	0.63 [0.48, 0.80]
GDPDEF	0.95 [0.92, 0.97]	1.81 [1.77, 1.84]	1.95 [1.86, 2.06]
Term spread	0.63 [0.62, 0.64]	0.93 [0.92, 0.93]	1.04 [1.01, 1.06]
EBP	0.78 [0.77, 0.79]	1.38 [1.36, 1.41]	0.41 [0.37, 0.44]

Notes: Entries are posterior medians with 68% credible bands in brackets over the retained draws, weighted by the volume weights. Rows index the structural equations, labeled by the variable on which each equation is normalized. Regime 1 equals one by normalization. Regime dates follow Table 1.

E Impulse Responses to the Demand, Supply, MP, and EBP Shocks

Shocks

Figure E.1 reports the impulse responses to the demand, supply, MP, and EBP shocks under the baseline identification.




Figure E.1: Impulse responses to the demand, supply, MP, and EBP shocks.

Note. Columns report the demand, supply, MP, and EBP shocks, and rows report the eight variables with their units. Each shock is scaled to one standard deviation. Under Assumption 2 the SOMA impact response to all four shocks is zero, and the one-percent-SOMA normalization of the balance-sheet figures does not apply. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. Each row shares one vertical range across the four shocks.

F The RFBS Shock with the EBP Held at Zero

Figure F.1 reports the responses of the other five variables in the counterfactual of Section 5.2.1.

Figure F.2 repeats Figure 8 with the penalized offset, which keeps the EBP response near zero and keeps all 10,000 draws.




Figure F.1: Responses of the other variables to the RFBS shock with the EBP held at zero.

Note. Responses to the RFBS shock, normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid blue lines and blue regions show the RFBS shock alone. Dashed green lines and green regions show the RFBS shock combined with the EBP shocks that keep the EBP response at zero in every month. Lines are posterior medians and shaded regions are 68% credible bands over the 9,826 draws in which those shocks stay finite. Vertical axes are in basis points for the one-year yield and the term spread, percent for SOMA and $\log(\text{Reserves} + \text{ON RRP})$, and percentage points for the ON RRP ratio.




Figure F.2: Responses to the RFBS shock with the EBP held near zero by a penalized offset.
Note. As [Figure 8](#), with the offsetting EBP shocks chosen by the penalized offset of [Section 5.2.1](#) and the bands computed over all 10,000 draws. Vertical axes are in percent for GDP and GDPDEF and in percentage points for the EBP.

G Robustness of the Balance-Sheet Impulse Responses

The figures below plot, for each robustness specification of Section 5.3, the responses of GDP, the deflator, and the EBP to the two shocks that Table 4 summarizes, one figure per specification, in the format of Figures 6 and 7.




Figure G.1: Robustness with Assumption 1 removed.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The model drops the four impact zeros of Assumption 1 and keeps the three SOMA zeros of Assumption 2. Everything else is as in the baseline.




Figure G.2: Robustness with Assumption 2 removed.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The model drops the three SOMA zeros of Assumption 2 and keeps the four impact zeros of Assumption 1. Everything else is as in the baseline.




Figure G.3: Robustness with the zero restrictions removed.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The model drops all seven zeros and keeps every sign and narrative restriction. Everything else is as in the baseline.




Figure G.4: Robustness with a single variance regime.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The model uses one variance regime and no latent scales, and both narrative anchors fall in that regime. Everything else is as in the baseline.




Figure G.5: Robustness with Gaussian errors.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The model removes the latent scales and keeps the four variance regimes. Everything else is as in the baseline.




Figure G.6: Robustness with regime breaks at the lower bound.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The regime breaks sit at the exits from and the return to the zero lower bound, 2015:M11, 2020:M3, and 2022:M2. Everything else is as in the baseline.




Figure G.7: Robustness with the TGA as an exogenous regressor.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The model adds the ratio of the TGA to the sum of reserves, ON RRP, and the TGA at lags zero to two. Everything else is as in the baseline.




Figure G.8: Robustness with Treasury issuance as exogenous regressors.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The model adds the bill share of marketable debt and the log change of coupon securities outstanding, each at lags zero to two. Everything else is as in the baseline.




Figure G.9: Robustness with monthly outcome variables.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The log of industrial production replaces GDP and the log of the CPI replaces the deflator, and every restriction moves to the new rows. Everything else is as in the baseline.




Figure G.10: Robustness with the two-year yield as the policy rate.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The two-year Treasury yield replaces the one-year yield, and every restriction on that row moves with it. Everything else is as in the baseline.




Figure G.11: Robustness with the GZ credit spread.

Note. Rows are the RFBS shock and the OFBS shock, each normalized to a one-percent impact decline in SOMA holdings and shown in the contractionary direction. Solid lines are posterior medians and shaded regions are 68% credible bands over the retained draws. The corporate bond spread of [Gilchrist and Zakrajšek \(2012\)](#) replaces the EBP and carries no restriction, as the EBP carries none. Everything else is as in the baseline.

H Variance Decomposition at Three Years

Table H.1 reports the shares of Figure 10 at 36 months with their 68% intervals.

Table H.1: Forecast-error variance decomposition.

	Share of forecast-error variance at 36 months, percent					
	Demand	Supply	MP	RFBS	OFBS	EBP
GDP	7.0 [3.0, 18.7]	13.3 [4.1, 30.5]	12.5 [3.9, 28.6]	5.7 [1.1, 15.4]	3.0 [0.6, 10.6]	4.9 [0.8, 16.3]
GDPDEF	8.0 [2.2, 24.8]	30.8 [13.0, 51.9]	10.2 [2.3, 27.4]	4.9 [0.8, 13.7]	7.4 [1.5, 19.4]	3.4 [0.6, 11.3]
EBP	5.8 [1.6, 17.1]	3.8 [1.0, 13.7]	6.1 [1.8, 19.0]	17.0 [4.9, 37.2]	3.8 [1.1, 12.1]	19.8 [6.1, 42.1]

Notes: Shares of the forecast-error variance in percent at a horizon of 36 months, under the average transmission, with the relative variances at their average of one across regimes. The first row of each pair is the posterior median and the second is the 68% credible interval, both over the retained draws. The six labeled shocks appear here, and the two unlabeled shocks account for the remainder.

I Historical Decompositions for All Shocks

Figure I.1 reports the decomposition of Figure 11 for the five variables that figure does not show.




Figure I.1: Historical decomposition by volatility regime, the other five variables.

Note. Contribution of each labeled shock to the accumulated forecast error of the variable over the window. Units are basis points for the one-year yield and the term spread, percent for SOMA holdings and log(Reserves + ON RRP), and percentage points for the ON RRP ratio. Bars are posterior medians and whiskers are 68% credible intervals, both over the retained draws. The windows are those of Figure 11. The two unlabeled shocks account for the remainder. In the size panel the RFBS contribution excludes zero in Regimes 1 to 3 and the OFBS contribution in Regime 3, over early QT2, and over the drain. In the ratio panel the OFBS contribution excludes zero in Regime 3, Regime 4, early QT2, and the drain. It is positive in Regime 3 and early QT2 and negative in Regime 4 and the drain.